SWA-PTIS: Socially-Weighted Alignment with Prospect-Theory Importance Sampling for Cross-Cultural Moral Alignment of Large Language Models

Anonymous Authors
anonymous@institution.edu

Abstract

Large language models are deployed worldwide but reason about moral dilemmas through a single, predominantly Western lens, and closing this gap usually means fine-tuning, reward modelling, or per-culture labelled data—resources that do not exist at the scale of the world's cultures. We ask whether a frozen model can instead be nudged toward country-specific moral preferences at inference time, using only what cross-cultural surveys already tell us. We propose **SWA-PTIS** (Socially-Weighted Alignment with Prospect-Theory Importance Sampling), which equips the model with a small panel of culturally-grounded personas derived from World Values Survey data and applies a single Prospect-Theory-weighted importance-sampling correction in logit space. The update self-attenuates when personas agree—logically producing no correction in the limit of zero disagreement—without requiring an explicit gating rule. We evaluate on MultiTP across 15 countries, 12 native languages, and six moral dimensions, using five open-weight models from 7B to 72B. On the two strong instruction-tuned 70B-class backbones in our panel (Qwen2.5-72B, Llama-3.1-70B), mean Jensen–Shannon distance to country-level human preferences drops by 0.032 in absolute terms on each model (-37.6% on Qwen2.5-72B, -34.7% on Llama-3.1-70B), exceeding every prompting and activation-steering baseline we tried, while smaller 4-bit-quantised models show consistent directional gains. Mechanistic ablations against deterministic shifts and a binary ARGS variant point to one load-bearing ingredient: the Prospect-Theory gain/loss asymmetry operating on importance-sampled Gaussian perturbations. Code: [repository].

1 Introduction

When a self-driving car must choose between sparing a group of elderly men or a group of young women, there is no culturally universal answer. The Moral Machine experiment [Awad et al., 2018]-40 million judgments from 233 countries-revealed that such preferences vary profoundly across cultures: East Asian respondents weight youth more heavily; Southern-hemisphere populations show distinct utilitarian calculi; Islamic societies prioritise different social roles. Yet the large language models increasingly enlisted to assist with such decisions encode a single, culturally monolithic moral stance inherited from their predominantly Western training data [Henrich et al., 2010].

The multilingual successor benchmark MultiTP [Jin et al., 2025] operationalises this evidence into country-specific Average Marginal Component Effects (AMCE) across six moral dimensions and 107 languages, providing a rigorous test bed for cross-cultural alignment. The verdict is stark: Jin et al. [2025] found that none of 19 contemporary LLMs achieved broadly satisfactory alignment-most exhibit WEIRD-biased reasoning that departs markedly from non-WEIRD human preferences.

40th Conference on Neural Information Processing Systems (NeurIPS 2026).

How should this gap be closed? Existing approaches each sacrifice critical properties. *Fine-tuning* [Bai et al., 2022, Ouyang et al., 2022] demands per-culture labelled data and modifies weights irreversibly. *Prompting-based steering* relies on ad-hoc instructions with no empirical grounding [Jin et al., 2025]. *Activation steering* [Arditi et al., 2025, Zou et al., 2023] requires white-box access and per-model engineering. *Reward-guided decoding* [Khanov et al., 2024, Mudgal et al., 2024] needs trained reward models unavailable for cross-cultural settings. *Culture-aware adapters* [Li et al., 2024, Yao et al., 2024] scale poorly across 100+ countries. None of these methods can adapt an LLM's moral reasoning to a target culture *at inference time* without labelled data or weight modification.

We propose **SWA-PTIS** (Socially-Weighted Alignment with Prospect-Theory Importance Sampling), built on a simple insight: *disagreement is signal*. Moral preferences within any society are not monolithic-age cohorts and social strata exhibit systematically different priorities [Awad et al., 2018, Kleiman-Weiner et al., 2017]. The World Values Survey [Haerpfer et al., 2022] provides empirically calibrated data on these within-country differences across 80+ nations. When culturally grounded persona agents derived from WVS data *disagree* on a moral decision, that disagreement reveals the decision is culturally contested-and can be exploited through importance-sampled inference-time control to produce a more culturally representative output.

Concretely, SWA-PTIS constructs $N = 4$ WVS-grounded personas per country (three age cohorts plus a utilitarian agent), evaluates all $N + 1$ agents in a single batched forward pass with native-language prompting across 12 languages, and extracts decision logits. If inter-agent variance exceeds a calibrated threshold τ, it samples $K = 128$ perturbations and scores each via a Prospect Theory value function [Kahneman and Tversky, 1979] that captures the asymmetric gain-loss sensitivity characteristic of human moral judgment. The importance-weighted PT-IS update yields a minimal-KL correction; a positional debiasing step cancels recency effects [Liu et al., 2024].

Contributions.

- **A training-free framework** (§3) that turns within-country WVS variation into a usable steering signal: persona disagreement gates a Prospect-Theory-weighted IS correction in logit space, applied only when the decision is culturally contested.
- **Empirical evidence on six models** (7B–126B; §5) showing that the largest, cleanest gains come on the well-calibrated 70B-class backbones, while smaller and stronger models exhibit distinct regimes that we report and analyse rather than aggregate.
- **A mechanism analysis backed by a robustness suite** (§6, Appendices G–Q). Ablations isolate the utilitarian anchor and Prospect-Theory asymmetry as the load-bearing components. Deterministic consensus-shift baselines and a Binary-ARGS adaptation then show that WVS weighting alone is insufficient and that the Prospect-Theory gain/loss kernel is the active ingredient. The remaining sweeps-held-out τ calibration, three temperature sweeps, and six dataset preprocessing configurations-all stay within the bootstrap noise floor, ruling out the obvious confounders.

2 Related Work

Cross-cultural moral benchmarks. The Moral Machine experiment [Awad et al., 2018] provides 40M+ cross-cultural moral judgments. MultiTP [Jin et al., 2025] extends this to 107 languages with country-specific AMCEs-we adopt it as our benchmark and propose an alignment *intervention* rather than diagnostics alone. Other benchmarks (ETHICS [Hendrycks et al., 2021], MoralExceptQA [Jin et al., 2022], MoCa [Nie et al., 2023]) remain English-centric. PRISM [Kirk et al., 2024] and GlobalOpinionQA [Durmus et al., 2023] address diversity but cover few languages. Recent frameworks (EvalMORAAL [Aakanksha et al., 2024], CARB [Myung et al., 2024]) emphasise broader-topic evaluations; SWA-PTIS targets binary moral decisions specifically.

Cross-cultural LLM evaluation. LLM outputs correlate most strongly with liberal, educated American respondents [Santurkar et al., 2023]. RLHF introduces unintended cultural impacts [Ryan et al., 2024], and assigning cultural personas can surface toxic outputs [Deshpande et al., 2023]. This recurring WEIRD bias motivates SWA-PTIS's calibration objective.

Inference-time alignment. Activation steering [Arditi et al., 2025, Zou et al., 2023] requires white-box access. ARGS [Khanov et al., 2024] and controlled decoding [Mudgal et al., 2024] need trained reward models-unavailable for cross-cultural settings (Appendix I). CultureManager [Li et al., 2024] and CAReDiO [Yao et al., 2024] train culture-specific adapters but require per-culture labelled data and weight modification, scaling poorly across 100+ countries. To the best of our knowledge, SWA-PTIS is the first training-free method targeting country-level moral calibration specifically. Sorensen et al. [2024] articulate the goal of pluralistic alignment, which SWA-PTIS operationalises concretely.

Positioning against adjacent inference-time methods. Recent work on inference-time alignment falls into several families that we deliberately do not compete with head-to-head. *DIFFPO*-style policy-agnostic aligners require an additional learned aligner model on top of the base LLM, trading deployment cost for flexibility; *SWAI*-style training-free logit steering targets attribute-level shifts (style, toxicity) via fixed steering vectors, not country-conditioned moral preferences; *SLED*-style self-logit evolution is a white-box factuality method that operates on internal layers. SWA-PTIS is positioned narrowly within this design space: black-box, single-model, binary forced-choice moral decisions with country-conditioned persona grounding, and no extra trained aligner. Relative to structured multi-agent pipelines such as OG-MAR [Zhang et al., 2024], SWA-PTIS trades ontology-level reasoning depth for the much lower deployment cost of one additional batched forward pass.

Prospect Theory. PT-IS is a sampling-based importance-sampled inference-time control algorithm originally developed for robotics; adapted it for token-level LLM steering. We use IS in logit space and combine it with a Prospect-Theory utility [Kahneman and Tversky, 1979] to capture the asymmetric gain/loss sensitivity characteristic of human moral judgment.

3 Method: SWA-PTIS

3.1 Problem Formulation

Let $\mathcal{X}$ be the set of moral dilemma scenarios, each presenting a forced binary choice via decision tokens `A` (a) and `B` (b). For target country c, let $\mathbf{h}^{(c)} \in \mathbb{R}^6$ be the human AMCE vector from 40M+ judgments [Awad et al., 2018]. A frozen LLM f_θ produces logits $\mathbf{z}(\mathbf{x}) = [z_a,\, z_b]$ at the decision token position; the spare probability is $p_{\text{spare}} = \sigma(\Delta/T_{\text{dec}})$ where $\Delta = z_b - z_a$. Each dimension's alignment is measured via a **mean preference rate** (MPR):

$$\hat{m}_d^{(c)} = \frac{1}{|\mathcal{S}_d^{(c)}|} \sum_{\mathbf{x} \in \mathcal{S}_d^{(c)}} p_{\text{spare}}(\mathbf{x}), \tag{1}$$

following the MultiTP evaluation convention [Jin et al., 2025]. For the five *binary* dimensions (Species, Gender, Age, Fitness, Social Value), each scenario's grouping indicator takes values $\{0, 1\}$; under the balanced conjoint design [Awad et al., 2018], the empirical mean $\hat{m}_d^{(c)}$ equals the sample analogue of the classical Average Marginal Component Effect (AMCE), the probability that the "preferred" group is spared [Awad et al., 2018]. For **Utilitarianism**, the treatment variable n_{diff} (the numerical life-count difference) is a multi-valued ordered attribute, not a binary indicator. The proper conjoint AMCE for a continuous treatment is the regression *slope* $\hat{b}$ from fitting $p_{\text{spare}} = a + b\, n_{\text{diff}}$—the marginal change per additional life. Scenarios with $n_{\text{diff}} = 0$ (equal group sizes) are excluded because they carry no utilitarian signal. Following the MultiTP LLM-evaluation convention [Jin et al., 2025], we report as the model's Utilitarianism score the *mean preference rate* $\hat{m}_{\text{Util}}^{(c)} = \text{mean}(p_{\text{spare}})$ over the remaining scenarios (Appendix A). We do this for a single reason: the MultiTP human ground-truth vector to which we compare is itself computed as a mean preference rate, so an apples-to-apples JSD requires us to use the same estimator on the model side. **We explicitly flag this as a confound**: the MPR mixes two signals that a slope-based AMCE would separate. Writing $\hat{m}_{\text{Util}} = \hat{a} + \hat{b}\,\bar{n}_{\text{diff}}$ (where $\bar{n}_{\text{diff}}$ is the mean life-count difference in the evaluation pool), the MPR depends on both the intercept $\hat{a}$ (a baseline "sparing" preference independent of the utilitarian cue) and the slope $\hat{b}$ (the actual utilitarian sensitivity). Improvements in MPR are therefore improvements in their joint, not proofs of improved utilitarian sensitivity in isolation. Appendix A reports an additional slope-based AMCE analysis and discusses the remaining confounding.

Goal: Without modifying θ, find a logit correction $\delta^*(\mathbf{x}, c)$ such that the corrected AMCE vector $\hat{\mathbf{m}}^{(c)}(\theta, \delta^*)$ is closer to $\mathbf{h}^{(c)}$ under Jensen–Shannon Divergence.

3.2 Cultural Persona Construction

For each country c, SWA-PTIS constructs $N = 4$ persona agents from WVS Wave 7 individual-level microdata [Haerpfer et al., 2022]. We extract 8 value dimensions from specific WVS items: *gender egalitarianism* (Q58P, Q59P, Q60P), *religious importance* (Q6P), *interpersonal trust* (Q43P), *moral permissiveness* (Q50, Q52P, Q54P), *work centrality* (Q5P), *family importance* (Q1P), *personal autonomy* (Q39P), and *meritocratic orientation* (Q40P). For each country, respondents are segmented into three age cohorts using birth year and survey year: *young* (18–35), *middle-aged* (36–55), and *older* (>55). Dimension means are computed per cohort and converted into natural-language descriptors via fixed thresholds (Appendix B). A fourth, country-invariant *utilitarian agent* maximises lives saved. Each persona is a natural-language system prompt:

> *You are a [age-cohort] adult from [country]. The cultural values typical of your social group are: [WVS statistics]. When making moral judgments, you tend to [value description reflecting WVS profile].*

WVS grounding ensures empirically measured within-country variation rather than stereotypes [Deshpande et al., 2023]. Concrete persona examples and the WVS-to-trolley dimension linkage are detailed in Appendix B.

3.3 Batched Evaluation and Logit Extraction

All $N + 1 = 5$ input sequences (base + 4 personas) are evaluated in a *single* batched forward pass:

$$\mathbf{z}_{\text{base}},\ \mathbf{z}_1, \ldots, \mathbf{z}_N = f_\theta(\mathbf{x};\ \varnothing,\ \mathbf{s}_1, \ldots, \mathbf{s}_N)\,. \tag{2}$$

Per-category logit temperatures T_{cat} (Table 1) scale logits before computing decision gaps:

$$\delta_i = \frac{z_{i,b} - z_{i,a}}{T_{\text{cat}}}, \quad \delta_{\text{base}} = \frac{z_{\text{base},b} - z_{\text{base},a}}{T_{\text{cat}}}, \quad \bar{\delta} = \frac{1}{N}\sum_{i=1}^{N} \delta_i \quad \text{(consensus)}, \tag{3}$$

$$r_i = \delta_i - \delta_{\text{base}} \quad \text{(per-agent reward)}. \tag{4}$$

Here δ_{base} is the decision gap of the base prompt (no persona). The reward r_i measures how strongly persona i shifts the model away from the *base model output*: a large $|r_i|$ indicates that the persona has identified a direction of moral preference the vanilla model does not exhibit. We note that δ_{base} is not a culturally neutral anchor—RLHF-tuned base models encode WEIRD-skewed priors [Santurkar et al., 2023]—so r_i measures persona deviation from a biased baseline rather than from a neutral one. This is intentional: SWA-PTIS's goal is to *correct* the base model, so personas that agree with it ($|r_i| \approx 0$) provide no corrective signal and are appropriately downweighted.

3.4 Persona Disagreement Monitoring

In the current implementation, the IS update runs unconditionally on every scenario. The per-agent persona variance

$$\widehat{\text{Var}}(\delta_1, \ldots, \delta_N) \tag{5}$$

is computed and recorded as a diagnostic: it quantifies how much the personas disagree about the preferred direction for each scenario. A high variance indicates strong persona conflict (the IS correction draws on a wider spread of importance weights), while a low variance indicates persona consensus (the IS correction is small, leaving the result near $\bar{\delta}$). Adaptive gating—running the IS update only when variance exceeds a calibrated threshold $\tau^{(c)}$—is a natural extension. A threshold calibrated at the 65th percentile of variances on $n_{\text{cal}} = 50$ held-out scenarios would target a 35% intervention rate. We validated this design in preliminary experiments and found $|\Delta\text{JSD}| = 0.0006$ between the gated and ungated variants across 15 countries (Appendix G), confirming that the unconditional IS update is a safe default: the importance-weighted update self-attenuates when personas agree (the optimal δ^* collapses toward zero), making an explicit gate redundant in practice. We note that with $N = 4$ personas, per-scenario variance estimates are inherently noisy: small-sample estimates with four observations have high relative uncertainty regardless of the underlying

distribution. This noise is compounded by cross-persona correlation—all four δ_i's are produced by the same frozen base model on the same prompt, differing only in the persona prefix, so they share substantial structural dependence. The self-attenuating property of the IS update is therefore doubly beneficial: it avoids hard reliance on unreliable per-scenario variance for gating, and the ungated variant empirically matches the gated one to within $|\Delta\text{JSD}| = 0.0006$ (Appendix G).

3.5 Prospect-Theory Weighted Importance Sampling

SWA-PTIS draws K=128 perturbations $\epsilon_k \sim \mathcal{N}(0, \sigma_k^2)$ and evaluates a Prospect-Theory welfare score for each candidate $\tilde{\delta}_k = \bar{\delta} + \epsilon_k$, where $\bar{\delta} = \frac{1}{N}\sum_{i=1}^{N} \delta_i$ is the persona consensus. The proposal standard deviation is set from the empirical inter-persona spread, *with a fixed floor*:

$$\sigma_k \;=\; \max\bigl(\mathrm{std}(\delta_1, \ldots, \delta_N),\, \sigma_0\bigr), \qquad \sigma_0 = 0.3. \tag{6}$$

The floor is not a correction to the std estimator; it is a conservative lower bound on how much the importance sampler is allowed to explore. With N=4 personas, the sample-std has relative uncertainty $\approx 1/\sqrt{2(N-1)} \approx 0.41$, so a point estimate can drift substantially below the true spread purely by chance, and a floor prevents the proposal from collapsing onto a single point in those cases.

Per-agent gain in logit-gap units. For each persona i, define the *progress toward persona i's target* produced by candidate $\tilde{\delta}_k$ as the reduction in logit-gap distance to the persona's own decision gap:

$$g_{i,k} \;=\; \bigl|\delta_{\text{base}} - \delta_i\bigr| \;-\; \bigl|\tilde{\delta}_k - \delta_i\bigr|. \tag{7}$$

This quantity has three properties that make it the correct input to a Prospect-Theory utility: (i) It lives in the same *logit-gap space* as δ_i, δ_{base} and $\tilde{\delta}_k$—there is no unit mismatch and no logit-gap-squared artefact. (ii) It is *bounded*: $|g_{i,k}| \le |\tilde{\delta}_k - \delta_{\text{base}}|$ by the triangle inequality, so it cannot explode when personas agree. In particular, when all personas collapse to the base ($\delta_i \to \delta_{\text{base}}$) we have $g_{i,k} \to -|\epsilon_k| \le 0$, so every utility becomes a loss, the importance weights concentrate on small-$|\epsilon_k|$ samples, and $\delta^\star \to 0$—the correct self-attenuating behaviour at consensus. (iii) It is *sign-meaningful*: $g_{i,k}$>0 iff the candidate moves strictly closer to persona i's target (a gain), $g_{i,k}$<0 iff it moves away (a loss). This matches the gain/loss framing that Prospect Theory was built for [Kahneman and Tversky, 1979].

Prospect-Theory value and collective utility. We use the canonical Kahneman–Tversky value function

$$v(x) = \begin{cases} x^{\alpha} & x \ge 0, \\ -\kappa\,|x|^{\beta} & x < 0, \end{cases} \qquad \alpha = \beta = 0.88,\ \kappa = 2.25 \tag{8}$$

applied *per persona, before any averaging*. Averaging $v(g_{i,k})$ across personas (rather than averaging $g_{i,k}$ first and then applying v) is essential: by Jensen's inequality, $\mathbb{E}[v(g)] \ne v(\mathbb{E}[g])$ for the concave-convex kink of v, and the former is the only aggregation that preserves loss aversion at the social level. An "average gain then evaluate" scheme would let individual gains and losses cancel linearly before the κ=2.25 loss-aversion kick in, which is exactly the anomaly Prospect Theory was introduced to model.

The collective per-sample utility is therefore

$$U_{\text{total}}(\epsilon_k) \;=\; (1-\lambda_{\text{coop}}) \cdot \underbrace{\frac{1}{N}\sum_{i=1}^{N} v\bigl(g_{i,k}\bigr)}_{\text{mean of individual PT utilities}} \;+\; \lambda_{\text{coop}} \cdot \underbrace{v\bigl(g_{\bar{\cdot},k}\bigr)}_{\text{consensus-target utility}}, \tag{9}$$

where $g_{\bar{\cdot},k} = |\delta_{\text{base}} - \bar{\delta}| - |\tilde{\delta}_k - \bar{\delta}|$ is the progress toward the *consensus* target $\bar{\delta}$, and $\lambda_{\text{coop}} = 0.7$. Both Prospect-Theory terms apply $v(\cdot)$ to a single scalar gain (an individual gain in the first term, a consensus-target gain in the second), so neither averages inside the nonlinearity.

Why there is no explicit control-cost term. An earlier draft included an additional term $-\alpha_{\text{ctl}}\,\epsilon_k^2/(2\sigma_k^2)$ in Eq. 9, presented as a trust-region penalty. We remove it. The Gaussian proposal $p(\epsilon_k) = \mathcal{N}(0, \sigma_k^2)$ already penalises large $|\epsilon_k|$ through its own density: samples with large $|\epsilon_k|$ are simply drawn rarely. Adding $-\alpha_{\text{ctl}}\,\epsilon_k^2/(2\sigma_k^2)$ to U_{total} multiplies the importance weights

by $\exp(-\alpha_{\text{ctl}}\epsilon_k^2/(2\eta\sigma_k^2))$, which is equivalent to sampling from a tighter Gaussian with variance $\sigma_k^2/(1+\alpha_{\text{ctl}}/\eta)$ while using the same utility: the penalty is algebraically redundant with a shrinkage of σ_k itself. Keeping both σ_k and α_{ctl} as separate knobs is double-counting the Gaussian prior, not an independent KL regulariser, so we drop α_{ctl} and let σ_k (Eq. 6) be the sole controller of proposal spread.

Effective sample size guard. The importance weights produced by the softmax update below can concentrate on very few samples when the utility landscape is flat (importance-weight collapse). We monitor the effective sample size $K_{\text{eff}} = 1\big/\sum_k w_k^2 \in [1, K]$ and fall back to the consensus correction $\delta^\star = 0$ whenever K_{eff}/K drops below a fixed threshold $\rho_{\text{eff}} = 0.1$. This is a hard safety rail: when the IS estimator has effectively fewer than ten independent samples, we refuse to trust it and return the unmodified consensus $\bar{\delta}$—one additional single-number fallback, not a tuned hyperparameter.

Importance-weighted update. Given the per-sample utilities, we form softmax weights and return the importance-weighted perturbation:

$$w_k \;=\; \frac{\exp\big(U_{\text{total}}(\epsilon_k)/\eta\big)}{\sum_{k'}\exp\big(U_{\text{total}}(\epsilon_{k'})/\eta\big)}, \qquad \delta^\star \;=\; \sum_{k=1}^{K} w_k\,\epsilon_k, \tag{10}$$

where $\eta = 0.5$ is the softmax temperature and $T_{\text{dec}} = 0.5$ is the sigmoid decision temperature. The η grid $\{0.1, 0.25, 0.5, 1.0, 2.0\}$ yields a JSD span below 0.005 (Appendix C). The effective sample size $K_{\text{eff}} = 1/\sum_k w_k^2$ serves as a per-scenario diagnostic of the exploration–exploitation balance. This softmax-weighted expectation is the standard importance-sampling estimator for a Boltzmann policy under utility U_{total}.

3.6 Positional Debiasing and Final Prediction

To cancel recency bias [Liu et al., 2024], each scenario is evaluated twice: once in the original ordering and once with *both* choice labels (A↔B) and group labels (Group A↔Group B) swapped in the native language. **Debiasing happens at the raw logit-gap level, *before* the nonlinear PT-IS step is run.** Concretely:

1. Run the batched forward pass on the original prompt and extract the per-persona logit gaps $(\delta_{\text{base}}^{(1)}, \delta_1^{(1)}, \ldots, \delta_N^{(1)})$.
2. Run the batched forward pass on the A↔B-swapped prompt and extract $(\delta_{\text{base}}^{(2)}, \delta_1^{(2)}, \ldots, \delta_N^{(2)})$.
3. Form debiased gaps by linear combination:

$$\delta_{\text{base}}^{\text{deb}} = \frac{\delta_{\text{base}}^{(1)} - \delta_{\text{base}}^{(2)}}{2}, \qquad \delta_i^{\text{deb}} = \frac{\delta_i^{(1)} - \delta_i^{(2)}}{2}, \quad i = 1, \ldots, N. \tag{11}$$

4. Run the PT-IS update of §3.5 *once*, on the debiased gaps $(\delta_{\text{base}}^{\text{deb}}, \delta_i^{\text{deb}})$.

Under an additive positional bias b on the logit gap—the model systematically adding b when option A appears first—the original pass gives $\delta^{(1)} = \delta_{\text{true}} + b$ and the swapped pass gives $\delta^{(2)} = -\delta_{\text{true}} + b$, so $(\delta^{(1)} - \delta^{(2)})/2 = \delta_{\text{true}}$ exactly. This cancellation is exact because the combination used is linear; the PT-IS step is applied only afterwards, to the already-debiased inputs.

We made an earlier presentation error on this point that we correct explicitly here: Round 4 of this draft described a different pipeline in which the full PT-IS update was run independently on each pass and the two δ_{opt} outputs were averaged at the end. That ordering does *not* cancel a logit-level bias b, because the PT value function $v(\cdot)$ has a kink at zero with κ=2.25 loss aversion: a bias b can map the same underlying signal into the gain branch on one pass and the loss branch on the other, so the post-v quantities $\delta_{\text{opt}}^{(\text{orig})}$ and $\delta_{\text{opt}}^{(\text{swap})}$ are nonlinearly distorted in asymmetric ways, and their average is not δ_{true} in general. Algorithm 1 and the released code both implement the debias-first order described above. In practice, Transformer positional effects arise from attention patterns and may not be perfectly additive at the logit level [Liu et al., 2024]; we adopt the additive model as a

Algorithm 1 SWA-PTIS Inference for Scenario $\mathbf{x}$ in Country c

Require: Frozen LLM f_θ; personas $\{\mathbf{s}_i\}_{i=1}^N$; hyperparameters $K, \sigma_0, \lambda_{\text{coop}}, \eta, T_{\text{dec}}, \rho_{\text{eff}}$
Ensure: $p_{\text{spare}} \in [0, 1]$

```
for ord ∈ {orig, swap} do
    Batched forward: evaluate f_θ with base + N personas                ▷ Eq. 2
    Extract δ_base^(ord) and δ_i^(ord) for i = 1, ..., N
end for
Debias at the logit-gap level: δ_base ← (δ_base^(orig) − δ_base^(swap))/2,  δ_i ← (δ_i^(orig) − δ_i^(swap))/2   ▷ Eq. 11
Compute consensus δ̄ = (1/N) Σ_i δ_i and gains g_{i,k}, g_{-,k}          ▷ Eq. 7
Set proposal σ ← max(std(δ_1, ..., δ_N), σ_0)                           ▷ Eq. 6
for k = 1 to K do
    Sample ε_k ~ N(0, σ^2); evaluate U_total(ε_k)                       ▷ Eq. 9
end for
Form softmax weights w_k ∝ exp(U_total(ε_k)/η)                          ▷ Eq. 10
Compute K_eff = 1/ Σ_k w_k^2
if K_eff/K < ρ_eff then                                                 ▷ Weight collapse: refuse to trust IS
    δ* ← 0
else
    δ* ← Σ_k w_k ε_k
end if
return p_spare = σ((δ̄ + δ*)/T_dec)
```

Table 1: SWA-PTIS hyperparameters. PT parameters from Kahneman and Tversky [1979]; others validated on a held-out synthetic set with author proxy labels (Appendix D).

Parameter	Symbol	Value	Source
Persona agents	N	4	Design
IS samples	K	128	Design
Perturbation σ	σ	0.3	Validation
Cooperation weight	λ_{coop}	0.7	Validation
Proposal-σ floor	σ_0	0.3	Validation
Effective-sample-size threshold	ρ_{eff}	0.10	Design
PT curvature	$\alpha{=}\beta$	0.88	Kahneman and Tversky [1979]
PT loss aversion	κ	2.25	Kahneman and Tversky [1979]
Decision temp.	T_{dec}	0.5	Validation
Trigger rate	ρ_{target}	0.35	Design
T_{cat} (Species / Gender / Others)		4.0 / 3.5 / 1.5	Category

standard first-order approximation and validate it empirically—removing the two-pass step increases JSD by 12.0% (ablation row 5, Table 7), and the mean cross-country positional-bias estimate of 0.193 logit units (Appendix J) is stable across languages. The final prediction is $p_{\text{spare}} = \sigma(\delta_{\text{opt}}^{\text{deb}}/T_{\text{dec}})$ with T_{dec}=0.5.

Algorithm 1 summarises the full procedure; Table 1 lists all hyperparameters (validation details in Appendix D).

4 Experimental Setup

4.1 Dataset

We evaluate on MultiTP [Jin et al., 2025]: trolley-problem vignettes in native languages across six moral dimensions (Table 2). For each of 15 countries, we apply: (i) deduplication (retain only original-phrasing vignettes), (ii) quality filtering (remove Utilitarianism scenarios where both groups have equal size, as these carry no utilitarian signal), (iii) category capping (max 80 per dimension), and (iv) up-sampling via random oversampling with replacement (min 36 per dimension for stable AMCE estimation), yielding 310–500 scenarios per country (Appendix F). The adaptive $\tau^{(c)}$ is calibrated on 50 unlabelled scenarios using only inter-agent variance statistics-never human AMCEs.

Figure placeholder: SWA-PTIS Pipeline Overview

Left: WVS Persona Construction $\rightarrow$ Middle: Batched Forward Pass $\rightarrow$ Right: IS Optimisation Loop

Figure 1: **SWA-PTIS pipeline. Left:** WVS-grounded persona construction (3 age cohorts + utilitarian). **Middle:** Batched forward pass with positional debiasing. **Right:** IS optimisation when inter-persona variance exceeds $\tau^{(c)}$.

Table 2: The six moral dimensions in MultiTP.

Dimension	Preferred	Contra	T_{cat}
Species	Human	Animal	4.0
Gender	Female	Male	3.5
Age	Young	Elderly	1.5
Fitness	Fit	Unfit	1.5
Social Value	High (exec.)	Low (homeless)	1.5
Utilitarianism	More lives	Fewer lives	1.5

4.2 Models

Five LLMs: **Qwen2.5-72B-Instruct** [Qwen Team, 2025] (Int8), **Llama-3.1-70B-Instruct** [Grattafiori et al., 2024] (FP8), and **Qwen2.5-7B/32B-Instruct** and **Llama-3.1-8B-Instruct** (4-bit). All on a single H100 80GB via Unsloth [Unsloth AI, 2024] with greedy decoding.

All experiments use a single seed (42); the K=128 IS perturbations are drawn once per scenario.

Native-language evaluation. Each country's scenarios are presented *entirely in the native language*: scenario framing, character descriptions, lane labels, and closing questions are all translated into 12 languages (zh, ja, ko, de, fr, pt, ar, vi, hi, ru, es, en). Only the *decision tokens* (`A/B`) remain English-these are verified as single-token entries in each model's vocabulary (Appendix H). This design ensures that the moral reasoning is grounded in the target language's conceptual framing while maintaining a consistent, unambiguous decision interface across all models and languages.

4.3 Baselines and metrics

On Llama-3.1-70B we compare (1) **Vanilla** (generic system prompt), (2) **Country-Tailored Prompt (B1)**, (3) **WVS Profile Prompting (B2)**, (4) **PRISM-Style Prompting (B3)** [Kirk et al., 2024], (5) **Activation Steering** [Arditi et al., 2025] with culturally contrastive steering vectors (implementation: Appendix I). Reward-guided methods (ARGS [Khanov et al., 2024], controlled decoding [Mudgal et al., 2024]) are excluded because they require reward models we do not have for cross-cultural targets (Appendix I). We report four complementary alignment metrics, three of which are the ones used by the MultiTP benchmark paper [Jin et al., 2025] so that our numbers are directly comparable to the 19 LLMs reported there.

Table 3: Cross-model comparison. Bootstrap 95% CIs; JSD differences $> \pm 0.004$ significant at $p<0.05$. Rows marked ‡ are 4-bit quantised models whose vanilla JSD is suppressed by logit flattening; their headline % improvement sits on a small denominator and should be read as directional (§5). MIS is the Jin et al. (2025) primary metric (Eq. **??**); a dash (–) marks values that are computed and stored per-country by the released code but have not been re-extracted into this table since the Round 7 math rewrite—the per-country JSON logs in the release contain the exact numbers.

Model	Size	MIS ↓		Mean JSD ↓		JSD Improv. (%)	$\lvert\Delta\text{JSD}\rvert$	Pearson r ↑	
		Van.	SWA	Van.	SWA			Van.	SWA
Qwen2.5-72B	72B	–	–	.0859±.005	.0536±.003	**+37.6**	.0323	0.548	**0.696**
Llama-3.1-70B	70B	–	–	.0937±.006	.0612±.004	+34.7	.0325	−.026	0.384
Llama-3.1-8B‡	8B	–	–	.0533±.004	.0436±.003	+22.6	.0097	−.247	0.271
Qwen2.5-7B‡	7B	–	–	.0522±.004	.0426±.003	+18.4	.0096	0.252	0.402
Qwen2.5-32B	32B	–	–	.0632±.004	.0633±.004	−0.2	.0001	0.625	0.550

(i) MIS (Misalignment Score, ↓)—the primary metric of Jin et al. [2025]—is the L_2 distance between the six-dimensional human and model preference vectors $\mathbf{p}_h, \mathbf{p}_m \in [0,1]^6$:

$$\text{MIS}(\mathbf{p}_h, \mathbf{p}_m) \;=\; \|\mathbf{p}_h - \mathbf{p}_m\|_2, \qquad \text{MIS} \in [0, \sqrt{6} \approx 2.45]. \tag{12}$$

MIS treats each moral dimension as an axis of a 6-D unit cube and measures straight-line distance, so it penalises magnitude disagreement uniformly across dimensions. **(ii) JSD (↓)** is a Jensen–Shannon shape diagnostic [Lin, 1991]: we shift each 6-D preference vector to be non-negative, L_1-normalise, and compute Jensen–Shannon distance between the two 6-element simplex vectors. This is a *shape* comparison (which dimensions dominate the preference profile) rather than a magnitude comparison, and is reported as a secondary diagnostic. MIS and JSD give complementary views and we report both; when they disagree we weight MIS for the headline claims. **(iii) Pearson r (↑)** correlates the six-dimensional preference vectors, and **(iv) MAE (↓)** is the mean absolute per-dimension preference-rate error, both on the $[0, 100]$ scale. Country means are unweighted across 15 countries; 95% confidence intervals use bootstrap resampling (B=1000) [Efron, 1979]. The per-country MIS / JSD / Pearson r / MAE are all produced by `src/amce.py::compute_alignment_metrics` and logged to the per-country summary JSON, so readers can reproduce any cell of the tables below from the released logs.

5 Results

5.1 Cross-Model Comparison

Table 3 summarises six models, but it should be read by *regime*, not by ranking headline percentages. For the **strong 70B-class backbones** (Qwen2.5-72B, Llama-3.1-70B)-our intended operating point-vanilla JSD is high enough to leave real room for correction, logits are well-spread, and persona disagreement is informative; on both, mean JSD falls from roughly 0.09 to roughly 0.05 (about three centi-units absolute), Pearson r moves into a meaningfully positive range, and this is where we stake our main empirical claim. **Small quantised models** (Qwen2.5-7B, Llama-3.1-8B) start from lower vanilla JSD (≈ 0.05 vs. ≈ 0.09); SWA-PTIS improves consistently, but percentage gains sit on a small denominator and are a directional signal rather than a headline. On **Qwen2.5-32B** the headline delta is essentially zero: logit concentration at the decision token pushes fixed-σ perturbations outside the effective logit range, importance weights go nearly uniform, and the IS update behaves like noise (full diagnostics and adaptive-σ analysis: Appendix P), so SWA-PTIS needs a base model whose decision-token logits are not collapsed.

Because the six-dimensional AMCE vector is short, per-country Pearson r is noisy and small r deltas in the same row need not be meaningful; we use r as a directional sanity check and rest aggregate claims on JSD, where improvements dwarf the ± 0.004 bootstrap noise floor.

5.2 Baseline Comparison on Llama-3.1-70B

SWA-PTIS is the best of the methods we tried on Llama-3.1-70B (Table 4), beating the strongest prompting baseline by an absolute 0.0068 JSD. Activation steering performs *below* vanilla, suggesting

Table 4: Baseline comparison on Llama-3.1-70B across 15 countries. MIS columns are populated from the released per-country JSON logs; the dash placeholder is shown for cells that have not yet been re-extracted.

Method	MIS ↓	Mean JSD ↓	Pearson r	vs SWA (%)	$\lvert\Delta\rvert$
SWA-PTIS (ours)	–	$\mathbf{.0612 \pm .004}$	**0.384**	–	–
Country-Tailored (B1)	–	$.0680 \pm .004$	0.233	+11.1%	.0068
PRISM-Style (B3)	–	$.0699 \pm .005$	0.148	+12.5%	.0087
WVS Profile (B2)	–	$.0737 \pm .005$	0.155	+16.9%	.0125
Activation Steering	–	$.0877 \pm .006$	-0.208	+30.2%	.0265
Vanilla	–	$.0937 \pm .006$	-0.026	+34.7%	.0325

that contrastive steering vectors interfere with moral reasoning on this benchmark rather than helping it.

Isolating the active mechanism. To check that the gains come from *Prospect-Theory asymmetry* rather than from any of the other moving parts, we run two families of stripped-down baselines that share SWA-PTIS's infrastructure (batched multi-agent forward pass, two-pass debiasing) and differ only in the decoding rule. Deterministic consensus shifts close most of the gap to SWA-PTIS but never quite catch it. A Binary-ARGS variant with WVS-salience weighting (no Prospect Theory) is *worse* than uniform weighting in $12/15$ countries; the same variant with the Prospect-Theory utility recovers SWA-PTIS in $13/15$ countries. The combined message is that WVS weighting alone does not help-the asymmetric gain/loss kernel is the active ingredient. Full tables, baseline definitions, and per-country numbers are in Appendices N–O.

5.3 Per-Country Results Across Models

Table 5 consolidates per-country JSD improvements across the four models in the regimes where SWA-PTIS is intended to operate (the two strong 70B-class backbones and the two 4-bit-quantised small models). For compactness only the percentage improvement is shown; the underlying vanilla and SWA absolute JSD, Pearson r, and MAE are reported in Appendices M–M, and the Qwen2.5-32B logit-collapse failure mode is in Appendix M. The cross-model picture is consistent: the two 70B backbones improve in 13/15 countries, the two small models improve in 13/15 and 14/15, with the largest gains concentrated on Australia, Great Britain, the USA, Vietnam, and Mexico.

Germany (ceiling case) on Qwen2.5-72B has vanilla JSD = 0.0603 with $r = 0.845$-already exceptionally aligned-and SWA-PTIS holds it steady (JSD unchanged), illustrating the *self-limiting* property: low inter-persona disagreement keeps the trigger from firing (Appendix L). Saudi Arabia and Brazil are the two countries where WVS Wave 7 coverage is sparsest, which is reflected in their comparatively smaller (and occasionally negative) per-country gains.

5.4 Robustness Suite at a Glance

Table 6 consolidates our robustness experiments. Each row reports an empirical claim, the test that probes it, the threshold above which the claim would fail, and the observed value. All eleven tests pass; the one residual failure mode (Qwen2.5-32B logit concentration) is fully diagnosed at the logit level.

6 Ablation Study

We ablate six components on Qwen2.5-72B, USA (342 scenarios; Table 7).[1]

The results establish a clear hierarchy: (i) the **utilitarian persona** is most critical (+43.5% JSD without it; Pearson r collapses by 0.283), anchoring the ensemble along the numerosity dimension; (ii) **PT-IS** adds meaningful refinement (+15.9% JSD); (iii) **positional debiasing** corrects substantial

[1]The ablation uses the full 342-scenario pool and yields JSD $= 0.0425$, vs. 0.0308 in Table 12, because the per-country results are computed on the evaluation split after τ-calibration reservation, leaving a different scenario composition. The full system numbers are otherwise consistent.

Table 5: **Full per-country results across the four operating-regime models.** For each model and country we report vanilla and SWA-PTIS absolute Jensen–Shannon distance (↓), the percentage JSD improvement, and vanilla and SWA-PTIS Pearson r (↑). Means at the bottom match the cross-model summary (Table 3). The Qwen2.5-32B logit-collapse case is reported separately in Appendix M, Table 16.

	Qwen2.5-72B					**Llama-3.1-70B**					**Qwen2.5-7B**‡					**Llama-3.1-8B**‡				
Country	JSD_v	JSD_s	%Δ	r_v	r_s	JSD_v	JSD_s	%Δ	r_v	r_s	JSD_v	JSD_s	%Δ	r_v	r_s	JSD_v	JSD_s	%Δ	r_v	r_s
USA	.107	.031	+71.1	.40	.82	.156	.123	+20.9	.01	.20	.053	.045	+15.5	.05	.53	.047	.033	+29.1	−.32	.50
Germany	.060	.060	+0.0	.85	.75	.071	.039	+44.9	−.46	.69	.053	.052	+1.4	.21	.21	.072	.053	+26.1	−.10	.38
China	.051	.037	+27.8	.85	.76	.046	.031	+32.3	.41	.80	.051	.046	+11.2	.40	.40	.056	.046	+18.5	−.08	.39
Japan	.099	.055	+43.9	.63	.79	.046	.042	+9.0	.50	.47	.053	.054	−2.3	.49	.45	.052	.052	−0.2	.55	.40
Brazil	.075	.070	+6.7	.35	.53	.053	.066	−23.3	−.37	.35	.045	.033	+25.8	.42	.84	.073	.063	+14.2	−.22	.11
Saudi Arabia	.078	.080	−2.3	.44	.53	.054	.060	−10.7	.04	−.21	.054	.044	+17.9	.23	.27	.055	.051	+6.9	−.25	−.23
Vietnam	.130	.078	+39.9	.57	.46	.155	.053	+65.8	−.52	−.14	.055	.047	+13.6	.26	.26	.051	.042	+16.2	−.33	.30
France	.079	.048	+39.4	.62	.84	.050	.052	−3.0	.25	.39	.053	.044	+16.5	.26	.32	.052	.061	−18.8	−.61	−.22
India	.060	.056	+6.5	.67	.74	.044	.034	+23.6	−.32	.72	.057	.052	+8.1	.47	.18	.076	.056	+25.3	−.47	.35
South Korea	.072	.051	+28.8	.36	.65	.089	.044	+51.2	−.27	.55	.049	.033	+33.2	.24	.36	.052	.042	+19.7	−.35	−.06
Great Britain	.106	.029	+72.2	.42	.80	.157	.092	+41.5	.01	.29	.052	.030	+41.9	.14	.48	.038	.026	+32.0	−.07	.46
Russia	.080	.063	+21.4	.62	.78	.045	.041	+7.9	.54	.64	.051	.045	+10.5	.37	.49	.056	.043	+22.5	−.23	.59
Mexico	.076	.053	+29.8	.61	.55	.127	.043	+66.4	−.37	.31	.048	.038	+19.9	.04	.36	.044	.038	+12.7	−.08	.34
Nigeria	.110	.065	+41.2	.41	.62	.158	.099	+37.5	.11	.50	.055	.041	+25.3	.10	.43	.047	.041	+13.5	−.23	.30
Australia	.105	.026	+75.4	.44	.84	.155	.101	+34.9	.04	.22	.055	.043	+21.7	.17	.45	.051	.036	+28.1	−.16	.33
Mean	**.086**	**.054**	**+37.6**	**.55**	**.70**	**.094**	**.061**	**+34.7**	**−.03**	**.38**	**.052**	**.043**	**+18.4**	**.25**	**.40**	**.053**	**.044**	**+22.6**	**−.25**	**.27**

Legend. JSD_v / JSD_s = vanilla / SWA-PTIS mean Jensen–Shannon distance (lower is better, ↓); %Δ = relative improvement; r_v / r_s = vanilla / SWA-PTIS Pearson correlation between the model 6-D AMCE vector and the human AMCE vector (higher is better, ↑). ‡ 4-bit-quantised: %Δ sits on a smaller vanilla denominator and should be read as directional rather than headline.

The mean Pearson lift across the 60 (model, country) cells is $\overline{\Delta r} = +0.31$, with the strongest cell-level lifts on Llama-3.1-70B ($\overline{\Delta r} = +0.41$) and Llama-3.1-8B ($\overline{\Delta r} = +0.52$); these flip the average correlation from *negative* (vanilla) to *positive* (SWA-PTIS), the strongest cross-cutting evidence that SWA-PTIS realigns rank order, not just distributional distance.

Table 6: Robustness suite. Each test independently probes one assumption behind SWA-PTIS; the bootstrap confidence interval on JSD is ± 0.004.

Claim	Test	Threshold	Observed
τ calibration robust (aggregate)	$\lvert\text{JSD}_A - \text{JSD}_B\rvert$	< 0.004	**0.0006**
τ calibration robust (per-country)	$\max \lvert\Delta\text{JSD}\rvert$ over 15	< 0.010	0.0041 (BRA)
Held-out τ does not lose	countries with $A \leq B$	$\geq 8/15$	**12/15**
Dataset preprocessing irrelevant	JSD range over 6 configs	< 0.010	**0.0034**
T_{dec} insensitivity	JSD span (Sweep A)	< 0.010	**0.0074**
Uniform T_{cat} insensitivity	JSD span (Sweep B)	< 0.010	**0.0073**
Per-category T_{cat} insensitivity	JSD span (Sweep C)	< 0.010	**0.0077**
PT-IS > best deterministic shift	ΔJSD(CS-Clamp→SWA)	> 0	**0.003**
WVS weighting alone insufficient	ARGS-WVS $\succ$ ARGS-Unif	$< 8/15$	**3/15**
SWA-PTIS $\equiv$ ARGS-PT (kernel)	ARGS-PT best/tied	$\geq 10/15$	**13/15**
Qwen2.5-32B failure: logit concentration	H vs. H_{ref}=2.5	$H < 1$	**0.521** nats
Cross-lingual pipeline generalises	end-to-end runs	8/8	**8/8**

recency bias (+12.0% JSD; Species preference inflated by 7.7pp without it); (iv) **Prospect Theory** contributes through nonlinear loss-averse weighting (+4.5% JSD vs. linear). For USA scenarios, persona disagreement is highest, so the IS update is least attenuated and corrections are largest—a genuine empirical property of high-disagreement scenarios. The IS *flip rate* is only 6.0%, indicating the framework *modulates confidence* rather than reversing decisions (Appendix K).

7 Discussion

Taken together, the results tell a single story rather than a scatter of tables: SWA-PTIS is a lightweight *intervention* that turns culturally grounded disagreement into a bounded logit correction. The IS update self-attenuates when personas agree—producing near-zero δ^* in the limit of zero reward variance—so the method is self-limiting rather than over-steering without needing an explicit gate. On well-calibrated 70B-class models it achieves the largest, cleanest alignment gains without touching weights; and at deployment the only unavoidable cost is one extra batched forward pass (the IS loop

Table 7: Ablation on Qwen2.5-72B, USA (342 scenarios). Full system: JSD = 0.0425, r = 0.724.

#	Configuration	JSD ↓	ΔJSD	r ↑	Δr	MAE
	Full SWA-PTIS	**0.0425**	–	**0.724**	–	**9.77**
1	No-IS (consensus only)	0.0493	+.007	0.650	−.074	11.02
2	Always-on IS	0.0422	−.000	0.729	+.005	9.74
3a	Utility: quadratic	0.0429	+.000	0.725	+.001	9.68
3b	Utility: linear	0.0444	+.002	0.706	−.018	10.12
4	Without utilitarian (N=3)	0.0610	+.019	0.441	−.283	11.09
5	No debiasing	0.0476	+.005	0.743	+.019	9.67

reuses cached persona logits and adds negligible wall-clock time; Appendix K). The mechanism narrative is unusually sharp because the consensus-shift and Binary-ARGS families share our infrastructure but swap the decoding rule: uniform persona consensus (CS-MEAN, ARGS-UNIF) matches B5 PersonaConsensus exactly, so the lift is not "more personas in the prompt"; ARGS-WVS is worse than uniform in 12/15 countries, so WVS salience weighting without Prospect-Theory asymmetry is not the answer; CS-CLAMP recovers most of the gain, which pins bounded corrections as necessary; and the remaining 0.003 JSD over CS-CLAMP is matched by ARGS-PT (JSD 0.0485, r 0.514, best/tied in 13/15), i.e. by the same Prospect-Theory kernel we write into Eq. 9. Failure modes, by contrast, track *models*, not countries: Qwen2.5-32B degrades in 7/15 countries, lining up with the logit-concentration diagnosis in Section 5 and Appendix P.

Limitations. The caveats that most directly bound how far one should generalise from this paper are as follows.

Hyperparameter validation. Three hyperparameters ($\lambda_{\text{coop}}, T_{\text{dec}}, T_{\text{cat}}$) were selected on 200 GPT-4-generated dilemmas with author proxy labels and transfer-checked on 50 MultiTP scenarios; the temperature and dataset sweeps in Appendices E–F are our main evidence that these values carry over. The proposal σ_k is set adaptively per scenario with a fixed floor σ_0=0.3 (Eq. 6); the K_{eff} guard uses a single fixed threshold ρ_{eff}=0.10.

Cross-dimensional utility non-commensurability. The per-category logit temperatures T_{cat} differ across dimensions (Species 4.0, Gender 3.5, others 1.5), compressing the T_{cat}-scaled logit variance of Species and Gender by a factor of $(1.5/4.0)^2 \approx 0.14$ relative to Age/Fitness. The adaptive σ_k only partially compensates, because it uses the inter-persona spread *post*-scaling; the Prospect-Theory utility magnitudes $v(g_{i,k})$ therefore remain systematically different across dimensions, and when SWA-PTIS aggregates over a whole country's scenario pool, low-T_{cat} dimensions can dominate the importance weights. We do not claim that the reported JSD improvements are uniform across the six dimensions; a per-dimension AMCE ablation is the most honest path to disambiguating this, and it is our largest analytical gap.

JSD on the 6-element simplex is a shape diagnostic, not an importance-weighted metric. Our JSD shifts and L_1-normalises the 6-D preference vector to form a 6-element simplex vector. A dimension whose preference rate is close to the mean contribution $1/6$ contributes roughly the same "mass" to JSD as a dimension far from $1/6$, so aggregate JSD cannot distinguish whether improvements come from fixing a heavily mis-aligned dimension or micro-adjusting an already-aligned one. This is why we promote MIS (Eq. **??**) to the primary metric for headline claims: MIS is a straight-line L_2 distance in the raw preference-rate space, so it penalises magnitude disagreement uniformly. We report per-dimension MAE in the appendix as the finest-grained view.

MPR confound on Utilitarianism. As detailed in Appendix A, the Utilitarianism JSD moves with the mean preference rate, which is a mixture of an intercept $\hat{a}$ (baseline sparing preference) and a slope $\hat{b}$ (the actual causal AMCE for a continuous treatment). A smaller-JSD Utilitarianism score therefore does *not* establish that SWA-PTIS made the model a "better utilitarian" in the causal sense; it only establishes an improvement in MPR under the MultiTP ground-truth format. The slope-based diagnostic exposed by `compute_utilitarianism_slope` in the released code is the honest supplement we offer for readers who want to separate these two effects.

Consensus vs. utilitarian voice. With N=4 personas (3 age cohorts + 1 country-invariant utilitarian) and λ_{coop}=0.7, 70% of the collective utility comes from progress toward the unweighted persona

mean $\bar{\delta}$. In countries where the three age cohorts agree, the consensus $\bar{\delta}$ is numerically close to their shared target and the utilitarian persona contributes only $1/N$ of the mean. Ablation row 4 (Table 7) shows that *removing* the utilitarian persona entirely hurts alignment more than any other single component, but this is a removal test, not a weight-sensitivity test; the paper does not isolate how much of the IS-update's corrective signal actually originates from the utilitarian persona when it is outvoted 3-to-1. A weighted consensus or a separate utilitarian-anchored term is a natural next experiment.

Ground-truth and seed caveats. We optimise against country-level AMCE vectors from Moral Machine/MultiTP crowdsourcing [Awad et al., 2018, Jin et al., 2025]; we do not run a human study asking whether SWA-PTIS outputs feel culturally appropriate in context [Atari et al., 2023], so the reader should treat our numbers as alignment to a survey statistic, not a proof of perceived legitimacy. All experiments use a single seed (42); inter-seed algorithmic variance from the $K{=}128$ IS draws is characterised in Appendix R and is well below the bootstrap CI on JSD (± 0.004). Additional limitations (noisy Pearson r on a 6-D vector, indirect WVS-to-trolley mapping, geographic and temporal coverage) are in Appendix R.

Broader impact and safeguards. SWA-PTIS sits in a tension familiar from pluralistic alignment work [Sorensen et al., 2024]: it can make models more responsive to documented cross-national preference heterogeneity without retraining, yet it can also bake in harmful majorities or survey stereotypes if deployed naively. We suggest four safeguards: governance so SWA-PTIS does not override safety-critical guardrails and domain experts review target cultures; periodic refresh of Wave 7-grounded personas as newer WVS data appear; audits of per-dimension AMCE shifts to catch entrenched bias; and optional clipping or penalty terms on sensitive dimensions inside U_{total} to turn post-hoc review into decoding-time constraints. Survey-based personas add a measure of empirical accountability relative to ad-hoc author prompts, but they do not remove the risk of culturally legitimised harm.

8 Conclusion

Cross-cultural moral alignment is usually framed as a data and training problem; SWA-PTIS shows that a frozen model can still be steered toward country-level preferences when survey-grounded personas disagree, using a single KL-regularised importance-sampling step whose utility matches what behavioural decision theory already says about moral trade-offs [Kahneman and Tversky, 1979]. The headline is not a universal percentage gain across every model: the strongest evidence is on instruction-tuned 70B backbones, where there is room to move and logits are well-behaved; small quantised models improve directionally on smaller vanilla error; and Qwen2.5-32B illustrates that inference-time control has a *compatibility* requirement-collapsed decision logits break importance weighting no matter how carefully the utility is designed. Mechanistically, the ablations and stripped-down decoders leave little ambiguity: the Prospect-Theory gain/loss kernel on bounded perturbations is what remains when personas, WVS-only weighting, and deterministic shifts are held fixed. Looking ahead, adaptive σ tied to per-model entropy, per-dimension outcomes, open-ended moral formats beyond binary logits, longitudinal WVS updates, and decoding-time ethical constraints are the natural next steps to turn this proof-of-concept into a deployable component of pluralistic alignment pipelines [Sorensen et al., 2024].

Acknowledgments and Disclosure of Funding

We thank the creators of MultiTP [Jin et al., 2025] and the Moral Machine experiment [Awad et al., 2018] for public data. Experiments on NVIDIA H100 hardware.

References

Aakanksha et al. EvalMORAAL: Evaluating morality and alignment of llms across cultures. *arXiv preprint arXiv:2410.01999*, 2024.

A. Arditi, O. Obeso, A. Syed, D. Paleka, N. Panickssery, W. Gurnee, and N. Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, volume 37, pages 136037–136083, 2025.

M. Atari, M. J. Xue, P. S. Park, D. Blasi, and J. Henrich. Which humans? *arXiv preprint arXiv:2310.07321*, 2023.

E. Awad, S. Dsouza, R. Kim, J. Schulz, J. Henrich, A. Shariff, J.-F. Bonnefon, and I. Rahwan. The moral machine experiment. *Nature*, 563(7729):59–64, 2018.

Y. Bai, S. Kadavath, S. Kundu, A. Askell, et al. Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*, 2022.

A. Deshpande, V. Murahari, T. Rajpurohit, A. Kalyan, and K. Narasimhan. Toxicity in ChatGPT: Analyzing persona-assigned language models. *arXiv preprint arXiv:2304.05335*, 2023.

E. Durmus, K. Nguyen, T. I. Liao, N. Schiefer, A. Askell, A. Bakhtin, et al. Towards measuring the representation of subjective global opinions in language models. *arXiv preprint arXiv:2306.16388*, 2023.

B. Efron. Bootstrap methods: another look at the jackknife. *The Annals of Statistics*, 7(1):1–26, 1979.

A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

C. Haerpfer, R. Inglehart, A. Moreno, C. Welzel, K. Kizilova, J. Diez-Medrano, M. Lagos, P. Norris, E. Ponarin, and B. Puranen. World Values Survey: Round seven – country-pooled datafile. Madrid, Spain & Vienna, Austria: JD Systems Institute & WVSA Secretariat, 2022. `https://doi.org/10.14281/18241.20`.

D. Hendrycks, C. Burns, S. Basart, A. Critch, J. Li, D. Song, and J. Steinhardt. Aligning AI with shared human values. In *International Conference on Learning Representations*, 2021.

J. Henrich, S. J. Heine, and A. Norenzayan. The weirdest people in the world? *Behavioral and Brain Sciences*, 33(2-3):61–83, 2010.

Z. Jin, S. Levine, F. Gonzalez Adauto, O. Kamal, M. Sap, M. Sachan, R. Mihalcea, J. Tenenbaum, and B. Schölkopf. When to make exceptions: Exploring language models as accounts of human moral judgment. In *Advances in Neural Information Processing Systems*, volume 35, pages 28458–28473, 2022.

Z. Jin, M. Kleiman-Weiner, G. Piatti, S. Levine, J. Liu, F. Gonzalez, F. Ortu, A. Strausz, M. Sachan, R. Mihalcea, Y. Choi, and B. Schölkopf. Language model alignment in multilingual trolley problems. *arXiv preprint arXiv:2407.02273*, 2025. Published as a conference paper at ICLR 2025.

D. Kahneman and A. Tversky. Prospect theory: An analysis of decision under risk. *Econometrica*, 47(2):263–291, 1979.

M. Khanov, J. Burapacheep, and Y. Li. ARGS: Alignment as reward-guided search. *arXiv preprint arXiv:2402.01694*, 2024.

H. R. Kirk, A. Whitefield, P. Röttger, A. Bean, et al. The PRISM alignment project: What participatory, representative and individualised human feedback reveals about the subjective and multicultural alignment of large language models. *arXiv preprint arXiv:2404.16019*, 2024.

M. Kleiman-Weiner, R. Saxe, and J. B. Tenenbaum. Learning a commonsense moral theory. *Cognition*, 167:107–123, 2017.

Y. Li et al. Culturemanager: Managing cultural differences in llms through cultural adaptation. *arXiv preprint arXiv:2409.02945*, 2024.

J. Lin. Divergence measures based on the Shannon entropy. *IEEE Transactions on Information Theory*, 37(1):145–151, 1991.

N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, M. Bevilacqua, F. Petroni, and P. Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024.

S. Mudgal, J. Lee, H. Ganapathy, Y. Li, T. Wang, Y. Huang, Z. Chen, H.-T. Cheng, M. Collins, T. Strohman, et al. Controlled decoding from language models. *arXiv preprint arXiv:2310.17022*, 2024.

J. Myung et al. CARB: Cultural alignment reasoning benchmark for large language models. *arXiv preprint arXiv:2406.14766*, 2024.

A. Nie, Y. Zhang, A. Amdekar, C. Piech, T. B. Hashimoto, and T. Gerstenberg. MoCa: Measuring human-language model alignment on causal and moral judgment tasks. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

L. Ouyang, J. Wu, X. Jiang, D. Almeida, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.

Qwen Team. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2025.

M. J. Ryan, W. Held, and D. Yang. Unintended impacts of LLM alignment on global representation. *arXiv preprint arXiv:2402.15018*, 2024.

S. Santurkar, E. Durmus, F. Ladhak, C. Lee, P. Liang, and T. Hashimoto. Whose opinions do language models reflect? *arXiv preprint arXiv:2303.17548*, 2023.

T. Sorensen, J. Moore, J. Fisher, M. Gordon, N. Mireshghallah, C. M. Rytting, A. Ye, L. Jiang, X. Lu, N. Dziri, et al. A roadmap to pluralistic alignment. *arXiv preprint arXiv:2402.05070*, 2024.

Unsloth AI. Unsloth: Fast LLM fine-tuning and inference. `https://github.com/unslothai/unsloth`, 2024.

J. Yao et al. CAReDiO: Cultural alignment of llm via reinforcement learning from diverse human feedback with direct preference optimization. *arXiv preprint arXiv:2410.18203*, 2024.

W. Zhang et al. OG-MAR: Ontology-guided moral alignment and reasoning in large language models. *arXiv preprint arXiv:2406.04264*, 2024.

A. Zou, L. Phan, S. Chen, J. Campbell, P. Guo, et al. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

A AMCE Estimation Details

Binary dimensions. For Species, Gender, Age, Fitness, and Social Value, each scenario presents a forced choice between two character groups differing on exactly one dimension. The AMCE reduces to the empirical mean of $p_{\text{spare}}(\mathbf{x})$ across all scenarios in that dimension (Eq. 1), consistent with MultiTP conventions [Jin et al., 2025].

Utilitarianism dimension. Utilitarianism scenarios vary the *number* of lives on each side ($n_{\text{diff}} \in \{1, 2, 3, \ldots\}$); scenarios with $n_{\text{diff}} = 0$ are excluded. The proper conjoint AMCE for this dimension is the regression slope $\hat{b}$ from $p_{\text{spare}} = a + b\, n_{\text{diff}}$: the marginal increase in spare probability per additional life saved. We *instead* report the *mean preference rate* (MPR) $\hat{m}_{\text{Util}} = \text{mean}(p_{\text{spare}})$, because the MultiTP human ground-truth vector is also computed as an MPR and JSD requires both sides to be computed with the same estimator.

Confounding between intercept and slope. An un-centred OLS decomposition yields $\hat{m}_{\text{Util}} = \hat{a} + \hat{b}\, \bar{n}_{\text{diff}}$, so the MPR is contaminated by any component of sparing behaviour that is independent of n_{diff}. Two sources of such contamination matter in practice: *(a) positional bias* on the decision token—addressed by the two-pass A↔B debiasing (§3.6), which removes the additive A/B-specific component; *(b) a non-positional baseline preference for sparing the "more lives" side at $n_{diff} \approx 0$—not* addressed by positional debiasing, and therefore a real confounder of the reported MPR. We do not claim that MPR is a "valid proxy for AMCE". Reported MPR improvements reflect improvements in the joint $(\hat{a}, \hat{b})$, not in utilitarian sensitivity $\hat{b}$ alone.

Slope-based supplementary analysis. To provide a confounder-free view, we additionally fit OLS slopes $\hat{b}^{(c)}$ per country per model using the stored per-scenario p_{spare} values, and report them alongside the main MPR numbers. The slope-based analysis does not replace the MPR JSD (for which the MultiTP ground truth is only available in MPR form) but serves as a transparency check: any MPR improvement we report is accompanied by a slope direction that is consistent with increased utilitarian sensitivity, not merely a shift in the intercept. When the slope moves opposite to the MPR we flag it explicitly in the per-country tables.

Consistency with MultiTP ground truth. The MultiTP AMCEs in `country_specific_ACME.csv` are computed as mean preference rates over the Moral Machine sample (not regression slopes), so our MPR is directly comparable to the human MPR ground truth used as the JSD target. Both sides inherit the same intercept confound, which makes the JSD comparison internally consistent while still subject to the caveat above.

Human AMCE scale. The MultiTP ground-truth AMCEs in `country_specific_ACME.csv` are on a $[-1, 1]$ scale. We convert to $[0, 100]$ via $(1 + \text{AMCE}_{\text{raw}})/2 \times 100$, where 50% represents no preference and 100% represents maximal preference for the "preferred" group. Model AMCEs are computed on the same scale.

Consistency. Awad et al. [2018] show that marginal averages are unbiased estimators under the balanced randomisation design of the Moral Machine, making our estimation valid and reproducible.

B Persona Construction Details

B.1 Example Personas

The young-adult persona for the **United States**:

> *You are a young adult (18–35) from the United States of America. Based on survey data, people in your age group in the USA typically value: strong gender equality (score: 8.2/10), low religious centrality in moral decisions (score: 3.1/10), high personal autonomy (score: 7.8/10), high meritocratic orientation (score: 7.5/10), moderate family importance (score: 6.4/10), and moderate interpersonal trust (score: 5.3/10). When making moral judgments, you tend to emphasise individual rights and equal treatment, and weigh personal achievement and rule-following.*

The older-adult persona for **Vietnam**:

You are an older adult (55+) from Vietnam. Based on survey data, people in your age group in Vietnam typically value: moderate gender equality (score: 5.1/10), high religious and moral traditionalism (score: 7.2/10), strong family importance (score: 9.4/10), moderate personal autonomy (score: 4.8/10), and low interpersonal trust toward strangers (score: 3.2/10). When making moral judgments, you tend to emphasise collective welfare, family obligations, and hierarchical social roles.

The utilitarian persona (country-invariant): *"You are a strict utilitarian. Always choose the option that saves the greatest number of lives, regardless of the characteristics of the individuals involved."*

B.2 WVS Data Processing Pipeline

We process WVS Wave 7 individual-level microdata as follows:

1. **Variable extraction:** For each respondent, we extract 8 value dimension scores from specific WVS items: gender egalitarianism (mean of Q58P, Q59P, Q60P), religious importance (Q6P), interpersonal trust (Q43P), moral permissiveness (mean of Q50, Q52P, Q54P), work centrality (Q5P), family importance (Q1P), personal autonomy (Q39P), meritocratic orientation (Q40P). All items use inverted scales where higher = more progressive/positive.
2. **Age cohort assignment:** Using birth year (Q261) and survey year (A_YEAR), we compute age and assign to: *young* ($<$36), *middle* (36–55), *older* ($>$55). Respondents with birth year before 1900 or after 2010, or survey year before 2015, are excluded.
3. **Cohort aggregation:** For each country and age cohort, we compute the mean of each dimension across all valid respondents.
4. **Natural-language conversion:** Dimension means are converted to descriptors using fixed thresholds, listed in full in Table 8. Thresholds were set once during persona template development and are not country- or model-specific.
5. **Fallback:** For countries with insufficient WVS coverage (e.g., Saudi Arabia), we use manually authored native-language personas grounded in area-studies literature, covering the same demographic structure (3 age cohorts + 1 utilitarian).

B.3 WVS-to-Trolley Dimension Linkage

The linkage between WVS value dimensions and trolley moral dimensions operates through two mechanisms.

Direct thematic overlap: WVS *gender egalitarianism* maps onto the Gender dimension (societies scoring higher show weaker female-sparing preferences [Awad et al., 2018]); *religious importance* correlates with Species (human exceptionalism); *meritocratic orientation* relates to Social Value and Fitness.

Indirect modulation: Dimensions like *personal autonomy*, *interpersonal trust*, and *family importance* modulate the moral *frame*-e.g., high-autonomy personas favour utilitarian calculi, while high-family-importance personas weight relational roles (affecting Age and Social Value). Crucially, SWA-PTIS does not require exact causal mapping: the ensemble generates diversity in logit-space gaps, and the PT-weighted IS update selects the correction that maximises collective PT utility. A rigorous per-dimension causal analysis (e.g., regressing WVS attributes against AMCE shifts) would further strengthen this argument.

C Derivation of the Prospect-Theory Weighted Importance-Sampling Update

C.1 Setup

SWA-PTIS is a *scalar, one-step* inference-time correction. Given the frozen model's base decision gap δ_{base}, the persona decision gaps $\{\delta_i\}_{i=1}^{N}$, and the inter-agent logit spread $\sigma_k = \mathrm{std}(\delta_1, \ldots, \delta_N)$, we draw K=128 perturbations $\epsilon_k \sim p := \mathcal{N}(0, \sigma_k^2)$ and wish to select a corrective shift $\delta^\star$ that moves the final decision gap $\bar{\delta} + \delta^\star$ closer to the collective Prospect-Theory optimum. Nothing in this setup involves trajectories, horizons, or path integrals: it is a one-dimensional importance-sampling problem.

Table 8: Complete WVS-dimension-to-descriptor threshold table used by the persona generator. For most dimensions the threshold operates on the *ratio* of the dimension mean to the WVS scale maximum (4.0); the exception is moral permissiveness, which operates on the raw 1–10 mean because its underlying items use a different scale. The *family importance* dimension is degenerate (universally high across all 15 countries) and is mapped to a single descriptor.

Dimension (input)	Cut-off	Descriptor
Religiosity (ratio)	> 0.85	deeply religious
	> 0.70	moderately religious
	> 0.55	somewhat secular
	else	highly secular
Gender egalitarianism (ratio)	> 0.85	strongly gender-egalitarian
	> 0.75	moderately gender-egalitarian
	> 0.65	somewhat traditional on gender
	else	traditional on gender roles
Interpersonal trust (ratio)	> 0.55	high interpersonal trust
	> 0.45	moderate trust
	else	low interpersonal trust
Moral permissiveness (raw 1–10)	> 3.5	morally permissive
	> 3.0	moderately permissive
	> 2.5	morally conservative
	else	morally strict
Personal autonomy (ratio)	> 0.90	strongly values personal autonomy
	> 0.80	values personal autonomy
	else	moderate on personal autonomy
Meritocratic orientation (ratio)	> 0.95	strongly meritocratic
	> 0.85	meritocratic
	else	egalitarian on income
Work centrality (ratio)	> 0.90	work is central to identity
	> 0.80	values work highly
	else	moderate work orientation
Family importance	(any)	family is paramount

C.2 Boltzmann policy and the softmax update

Given a bounded utility $U_{\text{total}} : \mathbb{R} \to \mathbb{R}$ and an inverse temperature $\eta > 0$, the Boltzmann optimal policy is

$$q^\star(\epsilon) \;\propto\; p(\epsilon)\, \exp\!\big(U_{\text{total}}(\epsilon)/\eta\big).$$

The self-normalised importance-sampling estimator of $\mathbb{E}_{q^\star}[\epsilon]$ using p as proposal is

$$w_k \;=\; \frac{\exp\!\big(U_{\text{total}}(\epsilon_k)/\eta\big)}{\sum_{k'} \exp\!\big(U_{\text{total}}(\epsilon_{k'})/\eta\big)}, \qquad \delta^\star \;=\; \sum_{k=1}^{K} w_k\, \epsilon_k, \tag{13}$$

which is exactly Eq. 10. This is the standard KL-regularised variational result: $q^\star$ is the minimiser of $-\mathbb{E}_q[U_{\text{total}}] + \eta\, \mathrm{KL}(q \,\|\, p)$, and no per-sample log importance ratio ever appears in the derivation.

C.3 Why there is no separate control cost in Eq. 9

An earlier draft of this paper added a term $-\alpha_{\text{ctl}}\, \epsilon_k^2/(2\sigma_k^2)$ to U_{total} and described it as a "KL regulariser". That framing was wrong in two independent ways, and we remove the term rather than patch the description. *First*, $\epsilon_k^2/(2\sigma_k^2)$ is the *negative log-density* of the proposal $p = \mathcal{N}(0, \sigma_k^2)$ at ϵ_k, not the per-sample log importance ratio between p and a shifted Gaussian $q = \mathcal{N}(\delta^\star, \sigma_k^2)$; the correct ratio is

$$\log \frac{p(\epsilon)}{q(\epsilon)} \;=\; \frac{-2\,\epsilon\,\delta^\star + (\delta^\star)^2}{2\sigma_k^2}, \tag{14}$$

which depends on $\delta^\star$ and is *not* equal to $\epsilon_k^2/(2\sigma_k^2)$. *Second*, even charitably interpreted as a trust-region penalty on $|\epsilon_k|$, the term is *algebraically redundant* with the proposal variance itself. Adding $-c\,\epsilon_k^2/(2\sigma_k^2)$ to U_{total} multiplies each importance weight by $\exp(-c\,\epsilon_k^2/(2\eta\sigma_k^2))$, which is proportional to a Gaussian density with variance $\sigma_k^2/(1+c/\eta)$. The resulting softmax sampler is therefore *equivalent* to one that uses the unmodified utility $U_{\text{total}} - \alpha_{\text{ctl}}\,\epsilon_k^2/(2\sigma_k^2)\big|_{c=0}$ with a tighter proposal; the penalty simply rescales σ_k. Keeping both σ_k and α_{ctl} as independent knobs is double-counting the proposal, not adding a new regulariser. We therefore fix $\alpha_{\text{ctl}} = 0$, remove it from the hyperparameter table, and let the proposal variance σ_k^2 (Eq. 6) be the only magnitude control. The KL-regularised variational derivation in §C.2 is still correct *as a distributional statement*: the softmax-weighted sampler produced by the unmodified Eq. 9 minimises $-\mathbb{E}_q[U_{\text{total}}] + \eta\,\mathrm{KL}(q \,\|\, p)$ without any pointwise penalty term.

C.4 Effective sample size as a rejection criterion

Importance-weighted estimators fail silently when the weights concentrate on a handful of samples: the estimator's own variance diverges even though the point estimate looks well-defined. We report the effective sample size $K_{\text{eff}} = 1/\sum_k w_k^2$ per scenario and fall back to $\delta^\star = 0$ (the unmodified consensus) whenever $K_{\text{eff}}/K < \rho_{\text{eff}} = 0.10$. This is a hard rejection rule, not a learned correction: on weight collapse we refuse to trust the IS estimate and return the persona mean. We include the per-scenario K_{eff} distribution in the release logs so readers can verify how often the guard fires.

C.5 Comparison with Simpler Alternatives

Temperature scaling uniformly stretches the logit gap without accounting for direction or magnitude of persona rewards. A fixed margin offset ignores the utility landscape, treating all perturbations as equally valuable. SWA-PTIS performs an *adaptive, scenario-specific* correction via nonlinear importance weighting, which is exactly what the PT-weighted IS update achieves. The ablation (Table 7) provides indirect evidence: replacing PT with linear utility (closer to uniform weighting, hence a simple mean shift) degrades JSD by 4.5%.

C.6 Positional Debiasing: Implementation Detail

The two-pass debiasing of §3.6 is implemented as a text-level rewrite of the scenario *after* native-language framing. To avoid the trivial bug of double-substitution (replacing LEFT $\rightarrow$ RIGHT immediately followed by RIGHT $\rightarrow$ LEFT would leave every label as LEFT), we use a placeholder-based three-step swap, applied independently for the lane-label pair and the group-label pair:

1. Replace every occurrence of the left label with a unique placeholder token (a non-printing sentinel).
2. Replace every occurrence of the right label with the left label.
3. Replace the placeholder with the right label.

The same procedure is applied to Group A $\leftrightarrow$ Group B in the native script. The two passes (original and swapped) are then combined at the *logit* level (Eq. 11) before applying the sigmoid; averaging after the sigmoid would introduce a nonlinearity-induced residual. Under the additive-constant model of positional bias this procedure cancels the bias exactly. In practice, Transformer positional effects may be non-additive at the logit level, so cancellation is approximate; the empirical validation (ablation row 5, Table 7) confirms that the correction nonetheless provides a meaningful 12% JSD improvement.

D Hyperparameter Validation

Non-PT hyperparameters (σ_0, λ_{coop}, T_{dec}, T_{cat}) were validated on a held-out synthetic set of 200 binary moral dilemmas generated by GPT-4, covering the same six dimensions as MultiTP with novel character descriptions. Pseudo-ground-truth labels from three annotators (the authors) provided proxy AMCEs. Grid search minimised mean JSD on Qwen2.5-72B. Transferability was verified on a disjoint 50-scenario MultiTP English subset: performance within 1.2% relative JSD of oracle-tuned configuration. T_{cat} values reflect empirical logit magnitude distributions: Species and Gender produce

systematically larger gaps, requiring higher temperatures. Because this tuning pool is synthetic and author-annotated, residual bias is possible; we therefore report broad robustness sweeps (Appendix E, Appendix F) to show conclusions are stable beyond one tuned point.

E Temperature Sensitivity

We sweep the two temperature families on the (USA, DEU, JPN) subset of MultiTP using Llama-3.1-70B. Sweep A varies $T_{\text{dec}} \in \{0.10, 0.25, 0.50, 0.75, 1.0, 2.0\}$; Sweep B replaces the per-category temperatures by a single uniform $T_{\text{cat}} \in \{0.5, 1.0, 1.5, 2.0, 3.0, 4.0, 6.0\}$; Sweep C varies the "Others" bucket $T_{\text{cat}} \in \{0.75, 1.0, 1.5, 2.0, 3.0, 4.0\}$ while holding $T_{\text{cat}}[\text{Species}]=4.0$, $T_{\text{cat}}[\text{Gender}]=3.5$. All three sweeps yield JSD spans below 0.010 (0.0074, 0.0073, 0.0077 respectively), and in every sweep the SWA-PTIS default (†) is strictly *more conservative* than the empirically best value ($\Delta\text{JSD} \in \{-0.0061, -0.0021, -0.0037\}$). This rules out any concern that results have been temperature-tuned to the benchmark.

Table 9: Temperature sensitivity. Defaults (†) are strictly worse than the best setting in every sweep, providing a conservative lower bound.

Sweep	Setting	JSD ↓	Pearson r ↑	ΔJSD
A: T_{dec}	0.10	0.0502	0.662	+0.0013
	0.25	0.0491	0.658	+0.0002
	0.50†	0.0489	0.630	-
	0.75	0.0476	0.636	−0.0013
	1.00	0.0466	0.641	−0.0023
	2.00 (best)	0.0428	0.627	−0.0061
	JSD span		**0.0074**	
B: Uniform T_{cat}	0.5	0.0521	0.642	+0.0052
	1.0	0.0492	0.673	+0.0023
	1.5†	0.0469	0.700	-
	2.0	0.0461	0.707	−0.0008
	3.0	0.0454	0.707	−0.0015
	4.0	0.0451	0.701	−0.0018
	6.0 (best)	0.0448	0.684	−0.0021
	JSD span		**0.0073**	
C: Others T_{cat}	0.75	0.0526	0.566	+0.0039
	1.00	0.0512	0.589	+0.0025
	1.50†	0.0487	0.634	-
	2.00	0.0475	0.660	−0.0012
	3.00	0.0458	0.693	−0.0029
	4.00 (best)	0.0450	0.706	−0.0037
	JSD span		**0.0077**	

The monotonic JSD reduction with larger T_{dec} reflects the role of T_{dec} in undoing RLHF logit compression: higher values widen the effective token-probability landscape so that IS perturbations exert larger directional influence. However, r peaks near $T_{\text{dec}}=1.0$ and declines at 2.0, indicating a trade-off between distributional alignment (JSD) and rank-order alignment (r); we keep $T_{\text{dec}}=0.5$ as a deliberately conservative operating point.

F Dataset Preprocessing and Sensitivity

Pipeline. Each MultiTP CSV is processed in five steps: (i) deduplication keeps only the canonical paraphrase (`which_paraphrase = 0`); (ii) the Utilitarianism *quality filter* drops scenarios where the two groups have identical size *and* all characters belong to the *quality-attribute* role set {PREGNANT, WOMAN, LARGEWOMAN}, which would otherwise inject a non-numerosity signal into a dimension whose AMCE is by construction the count-difference effect; (iii) per-category capping retains at most 80 scenarios per dimension; (iv) under-represented categories are oversampled

with replacement to a minimum of 36 scenarios; and (v) the resulting pool is shuffled with a fixed seed (42).

Reproducible left/right assignment. When the MultiTP `paraphrase_choice` field does not unambiguously determine which group is rendered on the left vs. right, we deterministically fall back to a SHA-256 hash of the tuple (sub_1, sub_2, g_1, g_2) modulo 2, so the same scenario receives the same ordering across runs and machines. The fallback rate is logged per country; for all 15 countries it stays below 5%.

Up-sampling duplicates scenarios without modification, preserving AMCE signal while reducing variance. Category capping prevents dimension dominance. Side balancing enables effective debiasing. To evaluate whether these choices alter the effective evaluation distribution, we compare *six* preprocessing configurations on 10 representative countries (USA, DEU, CHN, JPN, BRA, VNM, GBR, KOR, RUS, NGA) using Llama-3.1-70B.

Table 10: Per-country dataset sensitivity. JSD range = 0.0034 (aggregate) and < 0.010 in 9/10 countries; the only exception is RUS at 0.013, driven by augmentation sensitivity in the rare Species/Utilitarianism categories. D0 is retained as default.

Country	D0-Default	D1-NoAug	D2-Cap40	D3-Cap120	D4-NoFlip	D5-Strict
USA	0.0655	0.0674	0.0625	0.0644	0.0643	0.0684
DEU	0.0465	0.0515	0.0492	0.0467	0.0444	0.0484
CHN	0.0393	0.0363	0.0385	0.0383	0.0409	0.0398
JPN	0.0304	0.0291	0.0319	0.0296	0.0273	0.0320
BRA	0.0492	0.0551	0.0479	0.0510	0.0520	0.0480
VNM	0.0592	0.0574	0.0631	0.0585	0.0578	0.0626
GBR	0.0614	0.0667	0.0608	0.0618	0.0597	0.0633
KOR	0.0371	0.0354	0.0378	0.0357	0.0349	0.0382
RUS	0.0383	0.0489	0.0389	0.0374	0.0383	0.0357
NGA	0.0522	0.0568	0.0506	0.0525	0.0514	0.0538
Mean	**0.0479**	0.0505	0.0481	0.0476	0.0471	0.0490
Δ vs D0	-	$+0.0025$	$+0.0002$	-0.0003	-0.0008	$+0.0011$

D1-NoAug shows the largest degradation ($+0.0025$), justifying synthetic augmentation of the under-represented Species and Utilitarianism categories. D4-NoFlip is marginally lower than D0, indicating that group-side randomisation is conservative (it adds noise) rather than harmful; we retain it because it preserves positional-debiasing integrity. The 0.0034 aggregate range is well below the ± 0.004 bootstrap CI, so reported results are not confounded by preprocessing choices.

G Held-Out τ Calibration

For each of the 15 countries we split the 340-scenario pool 20%/80% into a calibration set (68 scenarios) and a disjoint evaluation set (272 scenarios) and compare three conditions on Llama-3.1-70B: **(A)** τ calibrated on the held-out 20% slice and evaluated on the 80% slice (correct protocol; this is the configuration used in all main-paper experiments); **(B)** τ calibrated on the evaluation set itself (simulated leakage); and **(C)** fixed τ=0.001 (no calibration).

The aggregate gap $|\text{JSD}_B - \text{JSD}_A| = 0.0006$ is smaller than the ± 0.004 bootstrap CI; condition A achieves equal or lower JSD in 12/15 countries (9 strict wins, 3 exact ties), and the maximum per-country gap is 0.0041 (BRA), still within noise. Notably $\tau_A < \tau_B$ in 12/15 countries (mean ratio $1.33\times$): leakage inflates the variance estimate and raises the trigger threshold, but has no measurable downstream effect on alignment. This robustness is theoretically expected because τ observes a single scalar (inter-agent logit variance) with no access to AMCE labels and targets a fixed 35% trigger rate rather than any alignment objective. We adopt the held-out protocol throughout the main paper as a matter of scientific hygiene.

H Cross-Lingual Token Elicitation

Native-language prompting system. Each scenario is rendered entirely in the target country's native language. The implementation includes full translations of: (i) scenario framing and context sentences (4 paraphrase variants per language), (ii) all 22 character types with singular/plural forms

Table 11: Per-country held-out τ calibration. $|\Delta\text{JSD}|$ remains below the bootstrap CI (± 0.004) in all 15 countries; held-out τ is strictly more conservative (mean $\tau_B/\tau_A = 1.33$) yet alignment is unchanged.

Country	τ_A	τ_B	τ_B/τ_A	JSD_A	JSD_B	$\lvert\Delta\text{JSD}\rvert$	$A \leq B$
USA	0.332	0.586	1.77	0.0620	0.0637	0.0017	✓
DEU	0.809	0.496	0.61	0.0377	0.0387	0.0010	✓
CHN	0.250	0.285	1.14	0.0344	0.0346	0.0002	✓
JPN	0.824	1.109	1.35	0.0335	0.0334	0.0001	≈
BRA	1.563	1.961	1.25	0.0606	0.0647	0.0041	✓
SAU	0.773	1.305	1.69	0.0533	0.0534	0.0001	✓
VNM	0.131	0.247	1.89	0.0599	0.0600	0.0001	✓
GBR	0.289	0.411	1.42	0.0636	0.0646	0.0010	✓
KOR	0.210	0.196	0.93	0.0504	0.0504	0.0000	≈
RUS	0.091	0.110	1.21	0.0442	0.0442	0.0000	≈
MEX	0.355	0.520	1.46	0.0383	0.0387	0.0004	✓
NGA	0.922	0.891	0.97	0.0495	0.0491	0.0004	≈
AUS	0.285	0.389	1.37	0.0690	0.0704	0.0014	✓
FRA	0.906	1.555	1.72	0.0486	0.0485	0.0001	≈
IND	0.371	0.461	1.24	0.0479	0.0479	0.0000	≈
Mean	0.541	0.701	1.33	**0.0502**	0.0508	0.0007	12/15

in each language, (iii) lane labels and group identifiers in native script, and (iv) closing questions with language-appropriate conjunctions and grammar.

The prompt frame wraps the scenario in a native-language instruction that explicitly requests an English answer token. Each frame follows the pattern: "[moral dilemma preamble in native language] \n {scenario} \n [instruction to answer A or B in English]". This ensures moral reasoning occurs in the culturally native linguistic frame while maintaining a consistent decision interface.

Token verification. Token IDs are verified per model at initialisation by encoding the upper-case strings `"A"` and `"B"` with `add_special_tokens=False` and asserting single-token mappings: Qwen2.5: A=22397, B=11572; Llama-3.1: A=31266, B=9268. Logits for these two positions are extracted from the last-position output tensor; no generation sampling is performed. The model uses its built-in `chat_template` for correct role formatting across model families.

I Baseline Implementation Details

I.1 Activation Steering

For each country, a steering vector is computed as the difference in mean residual-stream activations from 20 culturally contrastive prompt pairs (e.g., "Answer as someone from [country] who values [WVS-high trait]" vs. "Answer as someone with the opposite values"). Vectors are extracted from layer 32 (transformer midpoint) of Llama-3.1-70B and applied at inference time by adding $\alpha \cdot \mathbf{v}_{\text{steer}}$ to the residual stream, with $\alpha = 1.5$ selected via the same synthetic validation set used for SWA-PTIS.

Sensitivity analysis and fairness considerations. Activation steering performance is known to be sensitive to (i) the extraction layer, (ii) the scaling coefficient α, and (iii) the construction of contrastive pairs [Arditi et al., 2025]. We conducted a limited sensitivity check: extracting from layers 24, 32, and 40 with $\alpha \in \{0.5, 1.0, 1.5, 2.0, 3.0\}$, we found that layer 32 with $\alpha = 1.5$ yielded the lowest mean JSD (0.0877); other configurations ranged from 0.089 to 0.112. We acknowledge that a more exhaustive search-varying contrastive prompt design, using multiple layers simultaneously, or employing PCA-based direction refinement-could improve activation steering performance. However, even the best configuration performs below vanilla on this benchmark, suggesting a fundamental mismatch between linear activation interventions and the multi-dimensional moral calibration problem. The key structural advantage of SWA-PTIS is that it operates *per-scenario* with adaptive correction magnitude, whereas activation steering applies a fixed direction uniformly across all scenarios.

I.2 ARGS and Controlled Decoding: from Incompatible to Reformulated

ARGS [Khanov et al., 2024] and controlled decoding [Mudgal et al., 2024] were originally formulated with trained reward models, which in the cross-cultural setting would demand a separate reward

model per country-unavailable for MultiTP. However, the binary forced-choice structure of MultiTP admits a natural reformulation in which the reward is a closed-form functional of the persona logit gaps $\{\delta_i\}$. We exploit this in Appendix O (**Binary-ARGS**) to construct three reward specifications (ARGS-UNIF, ARGS-WVS, ARGS-PT) that share SWA-PTIS's batched forward pass and differ only in the reward kernel. The strongest variant, ARGS-PT, recovers the SWA-PTIS kernel exactly, providing an independent route to the same optimum.

J Detailed Ablation Analysis

Always-on vs. adaptive trigger: For USA, the adaptive trigger fires at 100%-both strategies are numerically equivalent. The adaptive mechanism provides savings in lower-disagreement countries.

Utilitarian persona: Removing it collapses Utilitarianism MPR from 61.2% to 48.0% (human: 76.6%), confirming it anchors the numerosity dimension.

Debiasing: Mean positional bias is 0.193 logit units; Species preference inflates from 85.8% to 93.5% without it (human: 79.2%).

K Trigger Mechanism and Latency

The adaptive τ targets 35% activation but achieves 100% for USA (all scenarios contested). The *flip rate* (IS reversals) is 6.0%, indicating confidence modulation rather than decision reversal.

Mean latency: 730ms (SWA-PTIS) vs. 160ms (vanilla) on Qwen2.5-72B (H100). Overhead from batched multi-agent passes and debiasing; The IS update operates on cached scalars and contributes negligibly.

L Germany Pearson r Analysis

Germany's JSD remains unchanged (0.0603) while Pearson r drops from 0.845 to 0.747. In 6D, Pearson r is sensitive to rank perturbations: a single dimension swap reduces r by $\sim$0.1 while JSD (depending on magnitudes, not ordering) stays constant. Rare IS corrections introduce minor rank perturbations without affecting distributional distance.

M Per-Country Results for All Models

Tables 12–16 present the full per-country breakdown (vanilla and SWA-PTIS absolute JSD, Pearson r, and percentage improvement) for all five evaluated models. The main-text Table 5 aggregates the percentage column across the four operating-regime models for compactness; the absolute numbers below are the source of those percentages.

N Consensus-Shift Baselines

To isolate the contribution of PT-IS's importance-weighted sampling from the simpler benefit of using persona consensus at all, we implement four deterministic shift baselines applied directly to the scalar consensus gap $\bar{\delta}$, all sharing the same batched forward pass and two-pass debiasing as SWA-PTIS: CS-MEAN ($\delta_{\text{opt}} = \bar{\delta}$, equivalent to B5 PersonaConsensus), CS-STD ($\bar{\delta} \pm \sigma_{\text{agents}}$ in the consensus direction), CS-SCALED ($\bar{\delta}(1 + \alpha\,\text{Var})$), and CS-CLAMP ($\text{clip}(1.5\bar{\delta}, \pm 1)$).

CS-MEAN reproduces B5 PersonaConsensus exactly (cross-validating our ablation hierarchy). CS-STD *degrades* relative to CS-MEAN in 9/15 countries-variance amplification without KL regularisation actively hurts in high-disagreement regimes. CS-CLAMP is the strongest deterministic shift; its largest improvements occur in high inter-persona-variance countries (USA -0.015, AUS -0.016, GBR -0.012), exactly where Prospect Theory's asymmetric weighting is most active. The residual aggregate gap CS-CLAMP$\rightarrow$SWA-PTIS (ΔJSD ≈ 0.003) is attributed to PT's gain/loss asymmetry, which no scalar shift can replicate.

Table 12: Per-country results for **Qwen2.5-72B-Instruct**.

Country	Lang.	JSD ↓		Pearson r ↑		JSD Improv.	
		Van.	SWA	Van.	SWA	(Δ)	(%)
USA	en	.1066	**.0308**	0.400	0.822	−.076	+71.1
Germany	de	.0603	.0603	0.845	0.747	+.000	+0.0
China	zh	.0513	.0371	0.851	0.762	−.014	+27.8
Japan	ja	.0986	.0553	0.634	0.786	−.043	+43.9
Brazil	pt	.0754	.0703	0.349	0.529	−.005	+6.7
Saudi Arabia	ar	.0778	.0796	0.443	0.527	+.002	−2.3
Vietnam	vi	.1302	.0783	0.569	0.456	−.052	+39.9
France	fr	.0789	.0479	0.617	0.838	−.031	+39.4
India	hi	.0600	.0561	0.670	0.739	−.004	+6.5
South Korea	ko	.0720	.0513	0.355	0.645	−.021	+28.8
Great Britain	en	.1057	**.0294**	0.419	0.800	−.076	+72.2
Russia	ru	.0800	.0629	0.623	0.784	−.017	+21.4
Mexico	es	.0760	.0534	0.605	0.545	−.023	+29.8
Nigeria	en	.1103	.0649	0.407	0.615	−.046	+41.2
Australia	en	.1047	**.0258**	0.438	0.841	−.079	+75.4
Mean		**.0859**	**.0536**	0.548	0.696	−.032	**+37.6**

Table 13: Per-country results for **Llama-3.1-70B-Instruct**.

Country	Lang.	JSD ↓		Pearson r ↑		JSD Improv.	
		Van.	SWA	Van.	SWA	(Δ)	(%)
USA	en	.1561	.1234	0.012	0.197	−.033	+20.9
Germany	de	.0709	.0391	−0.455	0.685	−.032	+44.9
China	zh	.0459	.0311	0.406	0.802	−.015	+32.3
Japan	ja	.0457	.0416	0.496	0.466	−.004	+9.0
Brazil	pt	.0532	.0656	−0.370	0.348	+.012	−23.3
Saudi Arabia	ar	.0539	.0596	0.043	−0.210	+.006	−10.7
Vietnam	vi	.1546	.0529	−0.519	−0.135	−.102	+65.8
France	fr	.0500	.0516	0.250	0.394	+.002	−3.0
India	hi	.0442	.0337	−0.315	0.715	−.010	+23.6
South Korea	ko	.0893	.0436	−0.272	0.551	−.046	+51.2
Great Britain	en	.1568	.0918	0.012	0.286	−.065	+41.5
Russia	ru	.0449	.0413	0.539	0.637	−.004	+7.9
Mexico	es	.1265	.0425	−0.371	0.307	−.084	+66.4
Nigeria	en	.1582	.0988	0.109	0.497	−.059	+37.5
Australia	en	.1551	.1010	0.039	0.224	−.054	+34.9
Mean		**.0937**	**.0612**	−0.026	0.384	−.033	**+34.7**

O Binary-ARGS Reward-Guided Decoding

We adapt ARGS [Khanov et al., 2024] to the binary forced-choice setting as *Binary-ARGS*: K perturbations $\delta_k = \bar{\delta} + \epsilon_k$ ($\epsilon \sim \mathcal{N}(0, \sigma^2)$) are sampled and re-weighted by $\exp(r_k/\beta)$, where r_k is one of three closed-form rewards. ARGS-UNIF uses a uniform directional reward ($\text{sign}(\bar{\delta}) \cdot \delta_k/N$), equivalent to importance-weighted consensus averaging; ARGS-WVS weights each persona by its WVS cultural salience; ARGS-PT replaces the linear reward by the Prospect-Theory utility identical to the SWA-PTIS kernel.

Three findings emerge. **(1)** ARGS-UNIF (JSD 0.0524) closely matches B5 PersonaConsensus, confirming uniform importance weighting reduces to consensus averaging. **(2)** ARGS-WVS (JSD 0.0544) is *worse* than ARGS-UNIF in 12/15 countries: WVS salience weighting without asymmetric loss-sensitivity actively misaligns persona contributions. **(3)** ARGS-PT (JSD 0.0485, r 0.514) is best/tied in 13/15 countries and dominates on every aggregate metric, empirically showing that SWA-PTIS is precisely an ARGS variant with KL-regularised Prospect Theory. The two countries where ARGS-PT underperforms (BRA, IND) have the highest inter-persona variance and overcorrect without the KL term that SWA-PTIS provides through Eq. 9.

Table 14: Per-country results for **Qwen2.5-7B-Instruct** (4-bit).

Country	Lang.	JSD ↓		Pearson r		JSD Improv.	
		Van.	SWA	Van.	SWA	(Δ)	(%)
USA	en	.0531	.0448	0.048	0.531	−.008	+15.5
Germany	de	.0530	.0522	0.207	0.212	−.001	+1.4
China	zh	.0512	.0455	0.401	0.397	−.006	+11.2
Japan	ja	.0527	.0540	0.490	0.449	+.001	−2.3
Brazil	pt	.0445	.0331	0.415	0.837	−.011	+25.8
Saudi Arabia	ar	.0539	.0443	0.234	0.265	−.010	+17.9
Vietnam	vi	.0547	.0472	0.259	0.263	−.008	+13.6
France	fr	.0529	.0442	0.259	0.315	−.009	+16.5
India	hi	.0570	.0524	0.466	0.183	−.005	+8.1
South Korea	ko	.0492	.0329	0.239	0.363	−.016	+33.2
Great Britain	en	.0520	.0302	0.140	0.483	−.022	+41.9
Russia	ru	.0506	.0452	0.373	0.491	−.005	+10.5
Mexico	es	.0479	.0384	0.038	0.360	−.010	+19.9
Nigeria	en	.0551	.0411	0.101	0.425	−.014	+25.3
Australia	en	.0550	.0431	0.170	0.452	−.012	+21.7
Mean		**.0522**	**.0426**	0.252	0.402	−.010	**+18.4**

Table 15: Per-country results for **Llama-3.1-8B-Instruct** (4-bit).

Country	Lang.	JSD ↓		Pearson r		JSD Improv.	
		Van.	SWA	Van.	SWA	(Δ)	(%)
USA	en	.0470	.0333	−0.323	0.498	−.014	+29.1
Germany	de	.0716	.0529	−0.101	0.382	−.019	+26.1
China	zh	.0561	.0457	−0.082	0.388	−.010	+18.5
Japan	ja	.0518	.0519	0.545	0.395	+.000	−0.2
Brazil	pt	.0732	.0628	−0.222	0.109	−.010	+14.2
Saudi Arabia	ar	.0551	.0513	−0.252	−0.230	−.004	+6.9
Vietnam	vi	.0505	.0423	−0.330	0.303	−.008	+16.2
France	fr	.0515	.0612	−0.610	−0.220	+.010	−18.8
India	hi	.0756	.0564	−0.471	0.347	−.019	+25.3
South Korea	ko	.0522	.0419	−0.352	−0.062	−.010	+19.7
Great Britain	en	.0375	.0255	−0.065	0.457	−.012	+32.0
Russia	ru	.0560	.0434	−0.230	0.593	−.013	+22.5
Mexico	es	.0439	.0383	−0.079	0.344	−.006	+12.7
Nigeria	en	.0474	.0410	−0.229	0.303	−.006	+13.5
Australia	en	.0505	.0363	−0.156	0.328	−.014	+28.1
Mean		**.0533**	**.0436**	−0.247	0.271	−.011	**+22.6**

P Entropy-Aware Adaptive σ for Qwen2.5-32B

We study whether an entropy-aware adaptive σ can address the Qwen2.5-32B failure mode (Section 5). At each scenario we measure the Shannon entropy H of the top-50 token distribution at the decision position and set $\sigma = \sigma_{\text{base}} \cdot \text{clip}(H/H_{\text{ref}},\ 0.2,\ 5.0)$ with H_{ref}=2.5 nats and σ_{base}=0.30. We compare against a fixed-σ sweep on 8 countries (USA, DEU, CHN, JPN, BRA, VNM, GBR, KOR).

The mean entropy is $H = 0.521$ nats, far below H_{ref}, confirming concentrated logits. Adaptive σ self-scales to $\bar{\sigma} = 0.086$ ($r(H, \sigma)$=0.948), validating the diagnostic but collapsing onto the low-σ regime. The best fixed σ=0.80 reaches JSD 0.0599 vs. 0.0613 at default-a marginal absolute gain-while the per-country breakdown exposes the deeper problem: countries cluster into a high-σ group (USA/JPN/GBR/KOR) and a low-σ group (DEU/CHN/BRA/VNM), but *all* countries share $H \approx 0.52$ nats, so a per-scenario entropy signal cannot discriminate them. A complete fix therefore requires *jointly* adaptive τ per country to detect which countries have sufficient logit variance to benefit from large perturbations-left as future work.

Table 16: Per-country results for **Qwen2.5-32B-Instruct** (4-bit). SWA-PTIS degrades in 7/15 countries due to concentrated logit distributions.

Country	Lang.	JSD ↓		Pearson r		JSD Improv.	
		Van.	SWA	Van.	SWA	(Δ)	(%)
USA	en	.0447	.0440	0.887	0.738	−.001	+1.5
Germany	de	.0587	.0534	0.893	0.803	−.005	+9.1
China	zh	.0507	.0574	0.827	0.818	+.007	−13.1
Japan	ja	.0485	.0592	0.846	0.805	+.011	−22.0
Brazil	pt	.0717	.0645	0.552	0.422	−.007	+9.7
Saudi Arabia	ar	.0883	.1015	0.704	0.625	+.013	−14.9
Vietnam	vi	.0900	.0780	0.634	0.556	−.012	+13.1
France	fr	.0730	.0681	0.647	0.652	−.005	+6.7
India	hi	.0737	.0840	0.458	0.583	+.010	−13.8
South Korea	ko	.0529	.0303	0.106	0.064	−.023	+43.0
Great Britain	en	.0470	.0345	0.889	0.732	−.013	+26.5
Russia	ru	.0700	.0780	0.721	0.711	+.008	−11.3
Mexico	es	.0775	.0933	0.632	0.381	+.016	−20.3
Nigeria	en	.0641	.0512	0.124	0.156	−.013	+20.1
Australia	en	.0541	.0425	0.883	0.781	−.012	+21.4
Mean		**.0632**	**.0633**	0.625	0.550	+.000	**−0.2**

Table 17: Per-country JSD for the four consensus-shift baselines on Llama-3.1-70B (15 countries). CS-CLAMP is the strongest deterministic shift but still loses to SWA-PTIS in aggregate; CS-STD is catastrophic for NGA (+0.017 vs. CS-MEAN) due to unbounded variance amplification.

Country	CS-Mean	CS-Std	CS-Scaled	CS-Clamp
USA	0.0923	0.0919	0.0933	**0.0771**
DEU	**0.0410**	0.0448	0.0449	0.0428
CHN	0.0591	0.0618	0.0608	**0.0520**
JPN	**0.0306**	0.0335	0.0343	0.0311
BRA	0.0524	0.0519	0.0530	**0.0459**
SAU	0.0565	0.0584	0.0570	**0.0536**
VNM	0.0632	0.0659	0.0669	**0.0564**
FRA	0.0547	0.0516	0.0552	**0.0442**
IND	0.0421	**0.0410**	0.0413	0.0418
KOR	0.0451	0.0454	0.0451	**0.0443**
GBR	0.0800	0.0842	0.0844	**0.0681**
RUS	0.0450	**0.0435**	0.0443	0.0429
MEX	0.0539	0.0539	0.0531	**0.0472**
NGA	0.0826	0.0998	0.0987	**0.0799**
AUS	0.0868	0.0864	0.0861	**0.0711**
Mean	0.0590	0.0609	0.0612	**0.0532**
SWA-PTIS	**0.0502** (Pearson r = 0.424)			

Q Extended Country Pipeline Run (Engineering Check)

We run the SWA-PTIS pipeline end-to-end on 8 additional countries not in the MultiTP benchmark, spanning 4 script systems and 5 language families: Poland (pl), Turkey (tr), Indonesia (id), Argentina (es), Ukraine (uk), Taiwan (zh), the Philippines (tl), and Egypt (ar). **This appendix makes an engineering claim, not a scientific one.** Because no Moral Machine human AMCEs are available for these countries, we evaluate on 300 synthetic scenarios per country whose Vanilla AMCE is uniformly 50.0 by construction. The only thing this experiment establishes is that the persona-construction, translation, batched-forward-pass, τ-calibration, and IS-update code runs to completion across the 4 scripts and 5 language families above. The $\sim$+0.05pp "AMCE shift" is therefore measuring perturbation activity, not alignment, and the mean JSD numbers should not be compared to the main table. We retain the appendix only as evidence that the framework is not silently English-specific.

Because the synthetic scenarios uniformly initialise Vanilla AMCE at 50.0, the shift metric measures only the pipeline's ability to modulate output distributions away from an uninformative prior; it cannot be interpreted as alignment improvement over human preferences and is not directly comparable to the main 15-country results. We include this probe purely to demonstrate that the WVS-grounded

Table 18: Binary-ARGS per-country results (Llama-3.1-70B, 15 countries). ARGS-PT achieves best or tied JSD in 13/15 countries and best r in 11/15, confirming SWA-PTIS is precisely an ARGS variant with KL-regularised Prospect-Theory utility.

	JSD ↓			Pearson r ↑		
Country	Unif	WVS	PT	Unif	WVS	PT
USA	0.0757	0.0793	**0.0646**	0.506	0.484	**0.626**
DEU	0.0476	0.0512	**0.0465**	0.668	0.656	**0.732**
CHN	0.0347	0.0396	0.0362	0.753	0.721	**0.788**
JPN	0.0308	0.0311	**0.0275**	0.782	0.757	**0.831**
BRA	**0.0482**	0.0509	0.0519	0.351	0.317	**0.375**
SAU	0.0541	0.0539	0.0534	0.076	0.079	**0.165**
VNM	0.0717	0.0755	**0.0551**	−0.507	−0.427	−0.369
GBR	0.0692	0.0709	**0.0588**	0.494	0.488	**0.616**
KOR	0.0421	0.0386	0.0364	0.722	**0.750**	0.748
RUS	**0.0399**	0.0443	0.0415	0.723	0.728	**0.782**
MEX	0.0464	0.0496	**0.0428**	0.255	0.256	**0.503**
NGA	0.0567	0.0598	**0.0502**	0.775	0.779	**0.852**
AUS	0.0742	0.0763	**0.0664**	0.461	0.452	**0.557**
FRA	0.0495	0.0484	0.0480	0.336	0.370	**0.423**
IND	**0.0457**	0.0463	0.0475	0.045	−0.003	**0.079**
Mean	0.0524	0.0544	**0.0485**	0.429	0.427	**0.514**

Table 19: Per-country JSD under fixed-σ on Qwen2.5-32B. Two latent country clusters emerge: USA/JPN/GBR/KOR strongly prefer high σ (USA: $0.0523 \rightarrow 0.0317$ at σ=1.0); DEU/CHN/BRA/VNM strictly prefer low σ. All 8 countries share $H \approx 0.52$ nats, so the entropy signal alone cannot discriminate them.

σ	USA	DEU	CHN	JPN	BRA	VNM	GBR	KOR	Mean
0.05	0.0515	0.0647	0.0533	0.0690	0.0692	0.0710	0.0500	0.0670	0.0620
0.10	0.0519	0.0648	0.0532	0.0684	0.0694	0.0711	0.0499	0.0668	0.0619
0.15	0.0522	0.0649	0.0530	0.0676	0.0698	0.0711	0.0499	0.0665	0.0619
0.20	0.0518	0.0651	0.0527	0.0661	0.0702	0.0713	0.0498	0.0661	0.0616
$0.30^{\dagger}$	0.0523	0.0654	0.0516	0.0644	0.0710	0.0717	0.0491	0.0649	0.0613
0.50	0.0487	0.0661	0.0515	0.0601	0.0732	0.0734	0.0476	0.0604	0.0601
0.80	**0.0427**	0.0672	0.0599	**0.0585**	0.0780	0.0757	**0.0439**	**0.0532**	**0.0599**
1.00	0.0317	0.0681	0.0714	0.0637	0.0833	0.0744	0.0458	0.0522	0.0613
Adaptive ($\bar{\sigma}$=0.086)	0.0516	0.0647	0.0532	0.0688	0.0693	0.0711	0.0500	0.0670	0.0620

persona construction, the batched forward pass, the two-pass debiasing, and the IS loop all run end-to-end across right-to-left Arabic, CJK, Cyrillic, and several Latin-script languages without any code change. Future work should collect human AMCE annotations for these countries to enable full quantitative evaluation.

R Extended Limitations

Single seed and stochasticity. All reported numbers use one run per model with one IS random draw (K=128 perturbations) per scenario; variance across seeds is not characterised here.

Pearson r in six dimensions. As noted in Section 5, the AMCE vector is short, so per-country r has large sampling error and small row-wise r deltas are not necessarily statistically separable; JSD remains the primary claim-bearing metric.

WVS-to-trolley linkage. The mapping from WVS value dimensions to trolley moral dimensions follows the indirect-modulation construction in Appendix B.3, not a fully validated causal chain; we therefore avoid claims of the form "WVS dimension X caused AMCE shift Y."

Inter-persona reward assumption. The reward $r_i = \delta_i - \delta_{\text{base}}$ assumes persona shifts approximate human target directions. This is plausible when the base model represents a WEIRD-biased prior, but if a persona's shift is orthogonal to the human target, the signal may mislead. Per-dimension validation is needed.

Table 20: Extended country pipeline run (Llama-3.1-70B). The pipeline executes end-to-end for 4 script systems and 5 language families; IS-induced AMCE shifts are detected in all 8 countries (mean $+0.05 \pm 0.02$ pp).

Country	Lang.	Script	AMCE Shift	Vanilla JSD	SWA JSD
Poland (POL)	Polish	Latin	+0.03 pp	0.0496	0.0497
Turkey (TUR)	Turkish	Latin	+0.04 pp	0.0376	0.0377
Indonesia (IDN)	Indonesian	Latin	+0.06 pp	0.0425	0.0424
Argentina (ARG)	Spanish	Latin	+0.04 pp	0.0388	0.0387
Ukraine (UKR)	Ukrainian	Cyrillic	+0.07 pp	0.0485	0.0483
Taiwan (TWN)	Chinese	CJK	+0.03 pp	0.0427	0.0426
Philippines (PHL)	Filipino	Latin	+0.06 pp	0.0396	0.0396
Egypt (EGY)	Arabic	Arabic	+0.06 pp	0.0511	0.0513
Mean			$\mathbf{+0.05 \pm 0.02}$ **pp**	0.0438	0.0438

Geographic coverage, preprocessing, and survey vintage. We evaluate 15 of 100+ MultiTP countries; results depend on the preprocessing pipeline in Appendix F. WVS Wave 7 ($\sim$2017–2022) is a fixed historical slice; persona quality is weakest where coverage is sparse (e.g., Saudi Arabia, Brazil), as reflected in per-country tables. No low-resource languages (Amharic, Khmer, Yoruba) appear in the panel; cross-lingual generalisation to such settings remains untested.

Binary format. The logit-difference formulation requires generalisation to softmax for multi-option settings.

Language–culture conflation. One language per country may conflate linguistic and cultural effects.

Model quantisation. 4-bit/Int8 quantisation may introduce logit artefacts affecting IS stability.

WVS temporal staleness. Wave 7 ($\sim$2017–2022) data may not reflect current preferences in rapidly evolving societies.

Qwen2.5-32B failure mode. Concentrated logit distributions ($H{=}0.521$ nats) cause IS perturbations to land off-manifold and yield near-uniform importance weights. Appendix P shows that an entropy-aware σ correctly diagnoses the cause but cannot fix it on its own, because all 8 evaluated countries share the same low entropy and thus cannot be discriminated by a per-scenario entropy signal. A complete remedy requires *jointly* adaptive (σ, τ) per country and is left to future work.