SWA-MPPI: Socially-Weighted Alignment with Model Predictive Path Integral Control for Cross-Cultural Moral Alignment of Large Language Models

Anonymous Authors
anonymous@institution.edu

Abstract

Large language models are deployed worldwide but reason about moral dilemmas through a single, predominantly Western lens. Closing this gap usually means fine-tuning, reward modelling, or per-culture labelled data—resources that do not exist at the scale of the world's cultures. We ask a narrower question: *can a frozen model be nudged toward country-specific moral preferences at inference time, using only what cross-cultural surveys already tell us?*

Our answer is **SWA-MPPI**. We give the model a small panel of personas grounded in World Values Survey data for the target country. When the personas *disagree* on a dilemma—and only then—we treat that disagreement as a signal that the decision is culturally contested, and apply a single, KL-regularised correction in logit space derived from Model Predictive Path Integral control with a Prospect-Theory utility. When the personas agree, the method does nothing.

We evaluate on MultiTP across 15 countries, 12 native languages, and six moral dimensions, on six open-weight models from 7B to 126B. On the two strong instruction-tuned 70B-class backbones in our panel (Qwen2.5-72B, Llama-3.1-70B), mean Jensen–Shannon distance to country-level human preferences drops by 0.032 in absolute terms on each model (-37.6% on Qwen2.5-72B, -34.7% on Llama-3.1-70B), exceeding every prompting and activation-steering baseline we tried. The remaining four models behave very differently: smaller 4-bit-quantised models show modest absolute gains ($\sim$0.01) on top of an already low vanilla JSD that we suspect partly reflects logit flattening rather than genuine alignment; Mistral-Large changes by less than the bootstrap noise floor; and Qwen2.5-32B is essentially unchanged. We deliberately do not aggregate these regimes into a single headline number. Mistral-Large illustrates a real failure mode: its JSD improves slightly but its 6-dimensional rank ordering does not, a dissociation we discuss honestly rather than average away. A residual failure on Qwen2.5-32B is fully traced to logit concentration at the decision token—a model-level diagnostic, not a silent error. Across countries, the largest gains cluster in Anglosphere populations, which we present as an unexpected pattern that *deserves explanation rather than as an explained insight*.

Held-out τ calibration, temperature and preprocessing sweeps, and mechanistic ablations against deterministic shifts and a binary ARGS variant together point to one load-bearing ingredient: the Prospect-Theory gain/loss asymmetry operating on KL-regularised perturbations. Code: `[repository]`.

40th Conference on Neural Information Processing Systems (NeurIPS 2026).

1 Introduction

When a self-driving car must choose between sparing a group of elderly men or a group of young women, there is no culturally universal answer. The Moral Machine experiment [Awad et al., 2018]-40 million judgments from 233 countries-revealed that such preferences vary profoundly across cultures: East Asian respondents weight youth more heavily; Southern-hemisphere populations show distinct utilitarian calculi; Islamic societies prioritise different social roles. Yet the large language models increasingly enlisted to assist with such decisions encode a single, culturally monolithic moral stance inherited from their predominantly Western training data [Henrich et al., 2010].

The multilingual successor benchmark MultiTP [Jin et al., 2025] operationalises this evidence into country-specific Average Marginal Component Effects (AMCE) across six moral dimensions and 107 languages, providing a rigorous test bed for cross-cultural alignment. The verdict is stark: Jin et al. [2025] found that none of 19 contemporary LLMs achieved broadly satisfactory alignment-most exhibit WEIRD-biased reasoning that departs markedly from non-WEIRD human preferences.

How should this gap be closed? Existing approaches each sacrifice critical properties. *Fine-tuning* [Bai et al., 2022, Ouyang et al., 2022] demands per-culture labelled data and modifies weights irreversibly. *Prompting-based steering* relies on ad-hoc instructions with no empirical grounding [Jin et al., 2025]. *Activation steering* [Arditi et al., 2025, Zou et al., 2023] requires white-box access and per-model engineering. *Reward-guided decoding* [Khanov et al., 2024, Mudgal et al., 2024] needs trained reward models unavailable for cross-cultural settings. *Culture-aware adapters* [Li et al., 2024, Yao et al., 2024] scale poorly across 100+ countries. None of these methods can adapt an LLM's moral reasoning to a target culture *at inference time* without labelled data or weight modification.

We propose **SWA-MPPI** (Socially-Weighted Alignment with Model Predictive Path Integral control), built on a simple insight: *disagreement is signal*. Moral preferences within any society are not monolithic-age cohorts and social strata exhibit systematically different priorities [Awad et al., 2018, Kleiman-Weiner et al., 2017]. The World Values Survey [Haerpfer et al., 2022] provides empirically calibrated data on these within-country differences across 80+ nations. When culturally grounded persona agents derived from WVS data *disagree* on a moral decision, that disagreement reveals the decision is culturally contested-and can be exploited through stochastic optimal control [Williams et al., 2018] to produce a more culturally representative output.

Concretely, SWA-MPPI constructs $N = 4$ WVS-grounded personas per country (three age cohorts plus a utilitarian agent), evaluates all $N + 1$ agents in a single batched forward pass with native-language prompting across 12 languages, and extracts decision logits. If inter-agent variance exceeds a calibrated threshold τ, it samples $K = 128$ perturbations and scores each via a Prospect Theory value function [Kahneman and Tversky, 1979] that captures the asymmetric gain-loss sensitivity characteristic of human moral judgment. The importance-weighted MPPI update yields a minimal-KL correction; a positional debiasing step cancels recency effects.

Contributions.

- **A training-free framework** (§3) that turns within-country WVS variation into a usable steering signal: persona disagreement gates a Prospect-Theory-weighted MPPI correction in logit space, applied only when the decision is culturally contested.

- **Empirical evidence on six models** (7B–126B; §5) showing that the largest, cleanest gains come on the well-calibrated 70B-class backbones, while smaller and stronger models exhibit distinct regimes that we report and analyse rather than aggregate.

- **A mechanism analysis backed by a robustness suite** (§6, Appendices G–Q). Ablations isolate the utilitarian anchor and Prospect-Theory asymmetry as the load-bearing components. Deterministic consensus-shift baselines and a Binary-ARGS adaptation then show that WVS weighting alone is insufficient and that the Prospect-Theory gain/loss kernel is the active ingredient. The remaining sweeps—held-out τ calibration, three temperature sweeps, and six dataset preprocessing configurations—all stay within the bootstrap noise floor, ruling out the obvious confounders.

2 Related Work

Cross-cultural moral benchmarks. The Moral Machine experiment [Awad et al., 2018] provides 40M+ cross-cultural moral judgments. MultiTP [Jin et al., 2025] extends this to 107 languages with country-specific AMCEs-we adopt it as our benchmark and propose an alignment *intervention* rather than diagnostics alone. Other benchmarks (ETHICS [Hendrycks et al., 2021], MoralExceptQA [Jin et al., 2022], MoCa [Nie et al., 2023]) remain English-centric. PRISM [Kirk et al., 2024] and GlobalOpinionQA [Durmus et al., 2023] address diversity but cover few languages. Recent frameworks (EvalMORAAL [Aakanksha et al., 2024], CARB [Myung et al., 2024]) emphasise broader-topic evaluations; SWA-MPPI targets binary moral decisions specifically.

Cross-cultural LLM evaluation. LLM outputs correlate most strongly with liberal, educated American respondents [Santurkar et al., 2023]. RLHF introduces unintended cultural impacts [Ryan et al., 2024], and assigning cultural personas can surface toxic outputs [Deshpande et al., 2023]. This recurring WEIRD bias motivates SWA-MPPI's calibration objective.

Inference-time alignment. Activation steering [Arditi et al., 2025, Zou et al., 2023] requires white-box access. ARGS [Khanov et al., 2024] and controlled decoding [Mudgal et al., 2024] need trained reward models-unavailable for cross-cultural settings (Appendix I). CultureManager [Li et al., 2024] and CAReDiO [Yao et al., 2024] train culture-specific adapters but require per-culture labelled data and weight modification, scaling poorly across 100+ countries. To the best of our knowledge, SWA-MPPI is the first training-free method targeting country-level moral calibration specifically. Sorensen et al. [2024] articulate the goal of pluralistic alignment, which SWA-MPPI operationalises concretely.

Positioning against adjacent inference-time methods. Recent work on inference-time alignment falls into several families that we deliberately do not compete with head-to-head. *DIFFPO*-style policy-agnostic aligners require an additional learned aligner model on top of the base LLM, trading deployment cost for flexibility; *SWAI*-style training-free logit steering targets attribute-level shifts (style, toxicity) via fixed steering vectors, not country-conditioned moral preferences; *SLED*-style self-logit evolution is a white-box factuality method that operates on internal layers. SWA-MPPI is positioned narrowly within this design space: black-box, single-model, binary forced-choice moral decisions with country-conditioned persona grounding, and no extra trained aligner. Relative to structured multi-agent pipelines such as OG-MAR [Zhang et al., 2024], SWA-MPPI trades ontology-level reasoning depth for the much lower deployment cost of one additional batched forward pass.

MPPI and Prospect Theory. MPPI [Williams et al., 2018] is a sampling-based stochastic optimal control algorithm originally developed for robotics; Huang et al. [2024] adapted it for token-level LLM steering. We use MPPI in logit space and combine it with a Prospect-Theory utility [Kahneman and Tversky, 1979] to capture the asymmetric gain/loss sensitivity characteristic of human moral judgment.

3 Method: SWA-MPPI

3.1 Problem Formulation

Let $\mathcal{X}$ be the set of moral dilemma scenarios, each presenting a forced binary choice via decision tokens `A` (a) and `B` (b). For target country c, let $\mathbf{h}^{(c)} \in \mathbb{R}^6$ be the human AMCE vector from 40M+ judgments [Awad et al., 2018]. A frozen LLM f_θ produces logits $\mathbf{z}(\mathbf{x}) = [z_\ell,\ z_r]$; the spare probability is $p_{\text{spare}} = \sigma(\Delta/T_{\text{dec}})$ where $\Delta = z_r - z_\ell$. Each dimension's AMCE is estimated as:

$$\hat{m}_d^{(c)} = \frac{1}{|\mathcal{S}_d^{(c)}|} \sum_{\mathbf{x} \in \mathcal{S}_d^{(c)}} p_{\text{spare}}(\mathbf{x}), \tag{1}$$

following the MultiTP evaluation convention. For the five binary dimensions (Species, Gender, Age, Fitness, Social Value), this reduces to the empirical mean of p_{spare}. For the Utilitarianism dimension, where the predictor is the *count difference* $|n_{\text{pref}} - n_{\text{non-pref}}|$ between groups, we fit $p_{\text{spare}} \sim a + b \cdot n_{\text{diff}}$ via OLS and evaluate at the mean observed n_{diff} (Appendix A).

Goal: Without modifying θ, find a logit correction $\delta^*(\mathbf{x}, c)$ such that the corrected AMCE vector $\hat{\mathbf{m}}^{(c)}(\theta, \delta^*)$ is closer to $\mathbf{h}^{(c)}$ under Jensen-Shannon Divergence.

3.2 Cultural Persona Construction

For each country c, SWA-MPPI constructs $N = 4$ persona agents from WVS Wave 7 individual-level microdata [Haerpfer et al., 2022]. We extract 8 value dimensions from specific WVS items: *gender egalitarianism* (Q58P, Q59P, Q60P), *religious importance* (Q6P), *interpersonal trust* (Q43P), *moral permissiveness* (Q50, Q52P, Q54P), *work centrality* (Q5P), *family importance* (Q1P), *personal autonomy* (Q39P), and *meritocratic orientation* (Q40P). For each country, respondents are segmented into three age cohorts using birth year and survey year: *young* (18–35), *middle-aged* (36–55), and *older* (>55). Dimension means are computed per cohort and converted into natural-language descriptors (e.g., a mean religiosity ratio >0.85 yields "deeply religious"). A fourth, country-invariant *utilitarian agent* maximises lives saved. Each persona is a natural-language system prompt:

> *You are a [age-cohort] adult from [country]. The cultural values typical of your social group are: [WVS statistics]. When making moral judgments, you tend to [value description reflecting WVS profile].*

WVS grounding ensures empirically measured within-country variation rather than stereotypes [Deshpande et al., 2023]. Concrete persona examples and the WVS-to-trolley dimension linkage are detailed in Appendix B.

3.3 Batched Evaluation and Logit Extraction

All $N + 1 = 5$ input sequences (base + 4 personas) are evaluated in a *single* batched forward pass:

$$\mathbf{z}_{\text{base}},\ \mathbf{z}_1, \ldots, \mathbf{z}_N = f_\theta(\mathbf{x};\ \varnothing,\ \mathbf{s}_1, \ldots, \mathbf{s}_N)\,. \tag{2}$$

Per-category logit temperatures T_{cat} (Table 1) scale logits before computing decision gaps:

$$\delta_i = \frac{z_{i,r} - z_{i,\ell}}{T_{\text{cat}}}, \quad \bar{\delta} = \frac{1}{N}\sum_{i=1}^{N} \delta_i \quad \text{(consensus)}, \tag{3}$$

$$r_i = \delta_i - \delta_{\text{base}} \quad \text{(per-agent reward)}. \tag{4}$$

The reward r_i measures how strongly persona i shifts the model *away from the culturally-neutral baseline*-a large $|r_i|$ indicates that the persona has identified a direction of moral preference the vanilla model does not exhibit.

3.4 Adaptive Conflict Detection

When personas agree, the consensus $\bar{\delta}$ suffices. When they disagree substantially, a controlled correction is warranted. SWA-MPPI uses a per-country threshold $\tau^{(c)}$ calibrated on $n_{\text{cal}} = 50$ unlabelled scenarios to target a 35% MPPI trigger rate:

$$\tau^{(c)} = \text{Percentile}_{65}\Big(\big\{\text{Var}(\delta_1^{(s)}, \ldots, \delta_N^{(s)})\big\}_{s=1}^{n_{\text{cal}}}\Big). \tag{5}$$

MPPI fires iff $\text{Var}(\delta_1, \ldots, \delta_N) \geq \tau^{(c)}$. Calibration uses only variance statistics and never observes human AMCEs. Throughout the main experiments, $\tau^{(c)}$ is calibrated on a 20% *held-out* slice and evaluated on the disjoint 80% slice; an extensive 15-country held-out vs. in-distribution comparison shows $|\Delta\text{JSD}| = 0.0006$ (below the bootstrap CI), confirming that τ does not overfit the evaluation distribution (Appendix G).

3.5 Prospect-Theory-Weighted MPPI Optimisation

When triggered, SWA-MPPI draws $K = 128$ perturbations $\epsilon_k \sim \mathcal{N}(0, \sigma^2)$ ($\sigma = 0.3$) and evaluates a collective welfare score for each $\tilde{\delta}_k = \bar{\delta} + \epsilon_k$.

Per-agent cooperative utility.

$$U_i(\epsilon_k) = (1 - \lambda_{\text{coop}})\, v(r_i\, \tilde{\delta}_k) \;+\; \lambda_{\text{coop}}\, v\Big(\bar{r}_{-i}\, \tilde{\delta}_k\Big), \tag{6}$$

with $\lambda_{\text{coop}} = 0.7$ and $v(\cdot)$ the Prospect Theory value function with $\alpha = \beta = 0.88$, $\kappa = 2.25$ [Kahneman and Tversky, 1979]. Both r_i and $\tilde{\delta}_k$ are scalar logit-gap quantities, and their product is a signed alignment score: it is positive when the candidate perturbation $\tilde{\delta}_k$ moves in the same direction as persona i's deviation r_i from the base model, and negative when it moves against it. The Prospect-Theory function v then maps this signed score into a utility that penalises misalignment more strongly than it rewards agreement (κ=2.25), giving a loss-averse generalisation of the directional inner product implicit in scalar consensus shifts. The natural alternative, $v(\tilde{\delta}_k - \delta_{\text{target}})$, requires committing to a single target δ_{target}—which is precisely what the multi-persona setup is designed to avoid.

$$v(x) = \begin{cases} x^{\alpha} & x \geq 0, \\ -\kappa\, |x|^{\beta} & x < 0. \end{cases} \tag{7}$$

Collective utility with KL regularisation.

$$U_{\text{total}}(\epsilon_k) = \frac{1}{N} \sum_{i=1}^{N} U_i(\epsilon_k) \;-\; \alpha_{\text{KL}}\, \frac{\epsilon_k^2}{2\sigma^2}, \tag{8}$$

where $\alpha_{\text{KL}} = 0.05$. The KL term derives from the free-energy objective of MPPI, preventing overcorrection in a theoretically grounded way (Appendix C).

MPPI update.

$$w_k = \frac{\exp\big(U_{\text{total}}(\epsilon_k)/\beta\big)}{\sum_{k'} \exp\big(U_{\text{total}}(\epsilon_{k'})/\beta\big)}, \qquad \delta^* = \sum_{k=1}^{K} w_k\, \epsilon_k, \tag{9}$$

where $\beta = T_{\text{dec}} = 0.5$ is the MPPI inverse temperature that governs the sharpness of importance weights. We set β equal to the decision temperature by design: this ensures that the MPPI optimisation operates on the same scale as the final sigmoid decision, so that a perturbation ϵ_k that shifts δ_{opt} by one unit produces a proportionate change in p_{spare}. Lower β would concentrate weights on fewer high-utility samples (sharper selection); higher β would spread weights more uniformly (softer averaging). At $\beta = 0.5$, the effective sample size $K_{\text{eff}} = 1/\sum_k w_k^2$ is typically 15–40 (out of $K = 128$), indicating moderate concentration that balances exploration with exploitation.

3.6 Positional Debiasing and Final Prediction

To cancel recency bias [Liu et al., 2024], each scenario is evaluated twice: once in the original ordering, once with *both* choice labels (A ↔ B) and group labels (Group A ↔ Group B) swapped in the native language. The debiased decision gap for each agent is $\delta_i = (\delta_i^{(\text{orig})} - \delta_i^{(\text{swap})})/2$. The final prediction is $p_{\text{spare}} = \sigma(\delta_{\text{opt}}/T_{\text{dec}})$ where $\delta_{\text{opt}} = \bar{\delta} + \delta^*$ and $T_{\text{dec}} = 0.5$.

Algorithm 1 summarises the full procedure; Table 1 lists all hyperparameters (validation details in Appendix D).

4 Experimental Setup

4.1 Dataset

We evaluate on MultiTP [Jin et al., 2025]: trolley-problem vignettes in native languages across six moral dimensions (Table 2). For each of 15 countries, we apply: (i) deduplication (retain only original-phrasing vignettes), (ii) quality filtering (remove Utilitarianism scenarios where both groups have equal size, as these carry no utilitarian signal), (iii) category capping (max 80 per dimension), and (iv) up-sampling via random oversampling with replacement (min 36 per dimension for stable AMCE estimation), yielding 310–500 scenarios per country (Appendix F). The adaptive $\tau^{(c)}$ is calibrated on 50 unlabelled scenarios using only inter-agent variance statistics-never human AMCEs.

Algorithm 1 SWA-MPPI Inference for Scenario $\mathbf{x}$ in Country c

Require: Frozen LLM f_θ; personas $\{\mathbf{s}_i\}_{i=1}^N$; threshold $\tau^{(c)}$; hyperparameters $K, \sigma, \lambda_{\text{coop}}, \alpha_{\text{KL}}, T_{\text{dec}}$
Ensure: $p_{\text{spare}} \in [0, 1]$

```
for order ∈ {orig, swap} do
    Batched forward: evaluate f_θ with base + N personas
    Extract logit pairs; compute δ_i^(ord)                    ▷ Eq. 2
end for
Average: δ_i ← (δ_i^(orig) − δ_i^(swap))/2                    ▷ Debiasing
Compute consensus δ̄, rewards r_i, variance ν                 ▷ Eqs. 3–4
if ν ≥ τ^(c) then                                             ▷ Conflict detected
    for k = 1 to K do
        Sample ε_k ~ N(0, σ²); compute U_total(ε_k)           ▷ Eqs. 6–8
    end for
    δ* = Σ_k w_k ε_k                                          ▷ Eq. 9
else  δ* ← 0
end if
return p_spare = σ((δ̄ + δ*)/T_dec)
```

Table 1: SWA-MPPI hyperparameters. PT parameters from Kahneman and Tversky [1979]; others validated on a held-out synthetic set with author proxy labels (Appendix D).

Parameter	Symbol	Value	Source
Persona agents	N	4	Design
MPPI samples	K	128	Design
Perturbation σ	σ	0.3	Validation
Cooperation weight	λ_{coop}	0.7	Validation
KL penalty	α_{KL}	0.05	Validation
PT curvature	$\alpha{=}\beta$	0.88	Kahneman and Tversky [1979]
PT loss aversion	κ	2.25	Kahneman and Tversky [1979]
Decision temp.	T_{dec}	0.5	Validation
Trigger rate	ρ_{target}	0.35	Design
T_{cat} (Species / Gender / Others)		4.0 / 3.5 / 1.5	Category

4.2 Models

Six LLMs: **Qwen2.5-72B-Instruct** [Qwen Team, 2025] (Int8), **Llama-3.1-70B-Instruct** [Grattafiori et al., 2024] (FP8), **Qwen2.5-7B/32B-Instruct** and **Llama-3.1-8B-Instruct** (4-bit), **Mistral-Large-2407** [Mistral AI Team, 2024] (4-bit). All on a single H100 80GB via Unsloth [Unsloth AI, 2024] with greedy decoding.

A note on stochasticity and seeds. We report this here because readers should know it before reading any number in §5: every result in this paper is from a *single* run per model with seed 42. Greedy decoding makes the per-scenario forward pass deterministic, but the $K{=}128$ MPPI perturbations $\epsilon_k \sim \mathcal{N}(0, \sigma^2)$ are drawn once from this same seed, so each reported δ^* is a single Monte-Carlo sample. We do not characterise across-seed variance of δ^* in this paper; the $K{=}128$ sample size keeps the per-scenario importance-weighted mean concentrated in expectation, but a proper across-seed study is left to future work and the reader should weigh all reported deltas accordingly.

Native-language evaluation. Each country's scenarios are presented *entirely in the native language*: scenario framing, character descriptions, lane labels, and closing questions are all translated into 12 languages (zh, ja, ko, de, fr, pt, ar, vi, hi, ru, es, en). Only the *decision tokens* (`A`/`B`) remain English-these are verified as single-token entries in each model's vocabulary (Appendix H). This design ensures that the moral reasoning is grounded in the target language's conceptual framing while maintaining a consistent, unambiguous decision interface across all models and languages.

Figure placeholder: SWA-MPPI Pipeline Overview

Left: WVS Persona Construction → Middle: Batched Forward Pass → Right: MPPI Optimisation Loop

Figure 1: **SWA-MPPI pipeline. Left:** WVS-grounded persona construction (3 age cohorts + utilitarian). **Middle:** Batched forward pass with positional debiasing. **Right:** MPPI optimisation when inter-persona variance exceeds $\tau^{(c)}$.

Table 2: The six moral dimensions in MultiTP.

Dimension	Preferred	Contra	T_{cat}
Species	Human	Animal	4.0
Gender	Female	Male	3.5
Age	Young	Elderly	1.5
Fitness	Fit	Unfit	1.5
Social Value	High (exec.)	Low (homeless)	1.5
Utilitarianism	More lives	Fewer lives	1.5

4.3 Baselines

On Llama-3.1-70B: (1) **Vanilla** (generic system prompt), (2) **Country-Tailored Prompt (B1)**, (3) **WVS Profile Prompting (B2)**, (4) **PRISM-Style Prompting (B3)** [Kirk et al., 2024], (5) **Activation Steering** [Arditi et al., 2025] using culturally contrastive steering vectors (details in Appendix I). Reward-guided methods (ARGS, controlled decoding) are excluded as they require reward models unavailable for cross-cultural settings (Appendix I).

4.4 Metrics

JSD (primary, ↓): Jensen-Shannon Distance between model and human AMCE vectors. **Pearson** r (↑): correlation of 6D AMCE vectors. **MAE** (↓): mean absolute AMCE error. All aggregated as unweighted means across 15 countries with bootstrap 95% CIs (1000 resamples).

5 Results

5.1 Cross-Model Comparison

Table 3 summarises performance across six models, but the right way to read it is by *regime*, not by ranking the percentages.

Strong 70B-class backbones (Qwen2.5-72B, Llama-3.1-70B). These are the intended operating point for SWA-MPPI: the vanilla JSD is high enough to leave real room for correction, the logits are well-spread, and persona disagreement is informative. On both models the mean JSD drops

Table 3: Cross-model comparison. Bootstrap 95% CIs; differences $>\pm 0.004$ significant at $p<0.05$. Rows marked ‡ are 4-bit-quantised models whose vanilla JSD is suppressed by logit flattening; their headline % improvement is amplified by a small denominator and should be read as directional only (§5). Mistral-Large's absolute change is within the bootstrap noise floor.

Model	Size	Mean JSD ↓		Improv. (%)	|ΔJSD|	Pearson r ↑	
		Van.	SWA			Van.	SWA
Qwen2.5-72B	72B	$.0859 \pm .005$	$.0536 \pm .003$	**+37.6**	.0323	0.548	**0.696**
Llama-3.1-70B	70B	$.0937 \pm .006$	$.0612 \pm .004$	+34.7	.0325	−.026	0.384
Llama-3.1-8B‡	8B	$.0533 \pm .004$	$.0436 \pm .003$	+22.6	.0097	−.247	0.271
Qwen2.5-7B‡	7B	$.0522 \pm .004$	$.0426 \pm .003$	+18.4	.0096	0.252	0.402
Mistral-Large	126B	$.0457 \pm .003$	$.0430 \pm .003$	+5.8	.0027	0.658	0.550
Qwen2.5-32B	32B	$.0632 \pm .004$	$.0633 \pm .004$	−0.2	.0001	0.625	0.550

from roughly 0.09 to roughly 0.05—an absolute reduction of about three centi-units—and Pearson r moves into a meaningfully positive range. This is the regime in which we make our main empirical claim.

Small quantised models (Qwen2.5-7B, Llama-3.1-8B). These models have *lower* vanilla JSD than their larger siblings (≈ 0.05 vs. ≈ 0.09). We are cautious about reading this as "the small models are already aligned": 4-bit quantisation flattens the decision-token logits toward chance, and many MultiTP AMCEs sit close to 0.5, so a near-uniform predictor can score well on JSD without expressing any genuine preference. SWA-MPPI does yield additional improvement on these models, but the percentage gain is amplified by an already small denominator and should be read as a directional signal, not a headline.

Mistral-Large. This is the most informative negative finding in the paper. Mean JSD improves only modestly, and Pearson r *decreases* for $11/15$ countries. We do not bury this in a footnote: it suggests that on a model whose vanilla logits are already comparatively well-aligned, the MPPI correction can compress the AMCE vector toward the population mean and improve distributional distance while degrading rank-order alignment along the 6-dimensional axes. We discuss this dissociation explicitly in §7.

Qwen2.5-32B. The headline delta is essentially zero. We trace this to logit concentration at the decision token: the fixed σ perturbations land outside the model's effective logit range, importance weights become near-uniform, and MPPI degenerates to random noise. The full logit-level diagnostic, the entropy-aware adaptive-σ analysis, and the per-country σ sweep are in Appendix P; the practical takeaway is that SWA-MPPI requires a base model whose decision-token logits are not collapsed.

The 6-dimensional AMCE vector is short, so a single-country Pearson r carries large sampling error and small r deltas in the same row are not necessarily meaningful. We therefore treat r as a directional sanity check on top of JSD; aggregate statements rely on JSD, where our improvements are large relative to the ± 0.004 bootstrap noise floor.

5.2 Baseline Comparison on Llama-3.1-70B

SWA-MPPI is the best of the methods we tried on Llama-3.1-70B (Table 4), beating the strongest prompting baseline by an absolute 0.0068 JSD. Activation steering performs *below* vanilla, suggesting that contrastive steering vectors interfere with moral reasoning on this benchmark rather than helping it.

Isolating the active mechanism. To check that the gains come from *Prospect-Theory asymmetry* rather than from any of the other moving parts, we run two families of stripped-down baselines that share SWA-MPPI's infrastructure (batched multi-agent forward pass, two-pass debiasing) and differ only in the decoding rule. Deterministic consensus shifts close most of the gap to SWA-MPPI but never quite catch it. A Binary-ARGS variant with WVS-salience weighting (no Prospect Theory) is *worse* than uniform weighting in $12/15$ countries; the same variant with the Prospect-Theory utility recovers SWA-MPPI in $13/15$ countries. The combined message is that WVS weighting alone does

Table 4: Baseline comparison on Llama-3.1-70B across 15 countries.

Method	Mean JSD ↓	Pearson r	vs SWA (%)	$\lvert\Delta\rvert$
SWA-MPPI (ours)	$\mathbf{.0612 \pm .004}$	**0.384**	–	–
Country-Tailored (B1)	$.0680 \pm .004$	0.233	+11.1%	.0068
PRISM-Style (B3)	$.0699 \pm .005$	0.148	+12.5%	.0087
WVS Profile (B2)	$.0737 \pm .005$	0.155	+16.9%	.0125
Activation Steering	$.0877 \pm .006$	−0.208	+30.2%	.0265
Vanilla	$.0937 \pm .006$	−0.026	+34.7%	.0325

Table 5: Per-country results for **Qwen2.5-72B-Instruct**.

Country	Lang.	JSD ↓		Pearson r ↑		JSD Improv.	
		Van.	SWA	Van.	SWA	(Δ)	(%)
USA	en	.1066	**.0308**	0.400	0.822	−.076	+71.1
Germany	de	.0603	.0603	0.845	0.747	+.000	+0.0
China	zh	.0513	.0371	0.851	0.762	−.014	+27.8
Japan	ja	.0986	.0553	0.634	0.786	−.043	+43.9
Brazil	pt	.0754	.0703	0.349	0.529	−.005	+6.7
Saudi Arabia	ar	.0778	.0796	0.443	0.527	+.002	−2.3
Vietnam	vi	.1302	.0783	0.569	0.456	−.052	+39.9
France	fr	.0789	.0479	0.617	0.838	−.031	+39.4
India	hi	.0600	.0561	0.670	0.739	−.004	+6.5
South Korea	ko	.0720	.0513	0.355	0.645	−.021	+28.8
Great Britain	en	.1057	**.0294**	0.419	0.800	−.076	+72.2
Russia	ru	.0800	.0629	0.623	0.784	−.017	+21.4
Mexico	es	.0760	.0534	0.605	0.545	−.023	+29.8
Nigeria	en	.1103	.0649	0.407	0.615	−.046	+41.2
Australia	en	.1047	**.0258**	0.438	0.841	−.079	+75.4
Mean		**.0859**	**.0536**	0.548	0.696	−.032	**+37.6**

not help—the asymmetric gain/loss kernel is the active ingredient. Full tables, baseline definitions, and per-country numbers are in Appendices N–O.

5.3 Per-Country Results: Qwen2.5-72B

Table 5: SWA-MPPI improves JSD in 13/15 countries. **An Anglosphere pattern, presented as an open question.** The largest gains occur for the three English-speaking countries in the panel (Australia, Great Britain, USA)—precisely the populations one would expect to be *best* served by an English-trained model. Their vanilla JSD is the highest of any country we evaluate. We do not have a clean causal explanation for this; we offer only an observation: whatever the strong 70B Qwen and Llama models have learned about "the average moral subject" lies further from the Moral Machine AMCEs of these three populations than from those of, e.g., Germany or China. We deliberately stop short of attributing this to any specific RLHF mechanism, since we have no intervention that isolates RLHF as the cause.

Germany (ceiling case): vanilla JSD = 0.0603 with $r = 0.845$-already exceptionally aligned. SWA-MPPI holds steady (JSD unchanged), demonstrating its *self-limiting* property: low inter-persona disagreement means the trigger rarely fires. The Pearson r decrease (0.845→0.747) reflects sensitivity to small rank perturbations in 6D, not distributional degradation (Appendix L).

Saudi Arabia (−2.3%) and Brazil (+6.7%) have limited WVS Wave 7 coverage, constraining persona quality.

5.4 Zero-Preprocessing MultiTP Evaluation

A natural reviewer question is what happens if SWA-MPPI is evaluated on MultiTP *exactly as released*, with none of our deduplication, capping, filtering, or upsampling. This is the cleanest test of whether our preprocessing choices systematically favour the method. We report this as a placeholder here; final numbers will be added in the camera-ready version.

Table 6: Zero-preprocessing MultiTP evaluation (placeholder; results to follow). Same models, same personas, same hyperparameters as the main table; the only change is that the input pool is the raw MultiTP release with no deduplication, no capping, no quality filtering, and no upsampling.

Model	Vanilla JSD	SWA-MPPI JSD	ΔJSD	Notes
Qwen2.5-72B	[TBD]	[TBD]	[TBD]	raw release
Llama-3.1-70B	[TBD]	[TBD]	[TBD]	raw release
Mistral-Large	[TBD]	[TBD]	[TBD]	raw release

We expect—based on the 0.0034 JSD spread already observed across our six preprocessing configurations (§F)—that the zero-preprocessing setting will move absolute numbers but leave the qualitative ordering intact. The point of this table is to make that expectation falsifiable.

5.5 Robustness Suite at a Glance

Table 7 consolidates our robustness experiments. Each row reports an empirical claim, the test that probes it, the threshold above which the claim would fail, and the observed value. All eleven tests pass; the one residual failure mode (Qwen2.5-32B logit concentration) is fully diagnosed at the logit level.

Table 7: Robustness suite. Each test independently probes one assumption behind SWA-MPPI; the bootstrap confidence interval on JSD is ± 0.004.

Claim	Test	Threshold	Observed
τ calibration robust (aggregate)	$\lvert \text{JSD}_A - \text{JSD}_B \rvert$	< 0.004	**0.0006**
τ calibration robust (per-country)	$\max \lvert \Delta\text{JSD} \rvert$ over 15	< 0.010	0.0041 (BRA)
Held-out τ does not lose	countries with $A \leq B$	$\geq 8/15$	**12/15**
Dataset preprocessing irrelevant	JSD range over 6 configs	< 0.010	**0.0034**
T_{dec} insensitivity	JSD span (Sweep A)	< 0.010	**0.0074**
Uniform T_{cat} insensitivity	JSD span (Sweep B)	< 0.010	**0.0073**
Per-category T_{cat} insensitivity	JSD span (Sweep C)	< 0.010	**0.0077**
MPPI $>$ best deterministic shift	ΔJSD(CS-Clamp→SWA)	> 0	**0.003**
WVS weighting alone insufficient	ARGS-WVS $\succ$ ARGS-Unif	$< 8/15$	**3/15**
SWA-MPPI $\equiv$ ARGS-PT (kernel)	ARGS-PT best/tied	$\geq 10/15$	**13/15**
Qwen2.5-32B failure: logit concentration	H vs. H_{ref}=2.5	$H < 1$	**0.521** nats
Cross-lingual pipeline generalises	end-to-end runs	8/8	**8/8**

6 Ablation Study

We ablate six components on Qwen2.5-72B, USA (342 scenarios; Table 8).[1]

The results establish a clear hierarchy: (i) the **utilitarian persona** is most critical (+43.5% JSD without it; Pearson r collapses by 0.283), anchoring the ensemble along the numerosity dimension; (ii) **MPPI** adds meaningful refinement (+15.9% JSD); (iii) **positional debiasing** corrects substantial recency bias (+12.0% JSD; Species preference inflated by 7.7pp without it); (iv) **Prospect Theory** contributes through nonlinear loss-averse weighting (+4.5% JSD vs. linear). The adaptive trigger fires at 100% for USA-a genuine empirical property of high-disagreement scenarios, not a calibration failure. The MPPI *flip rate* is only 6.0%, indicating the framework *modulates confidence* rather than reversing decisions (Appendix K).

7 Discussion

Key findings. SWA-MPPI offers three properties simultaneously: (i) meaningful alignment gains on well-calibrated 70B models without weight modification, (ii) graceful self-limiting when intervention

[1]The ablation uses the full 342-scenario pool and yields JSD $= 0.0425$, vs. 0.0308 in Table 5, because the per-country results are computed on the evaluation split after τ-calibration reservation, leaving a different scenario composition. The full system numbers are otherwise consistent.

Table 8: Ablation on Qwen2.5-72B, USA (342 scenarios). Full system: JSD = 0.0425, r = 0.724.

#	Configuration	JSD ↓	ΔJSD	r ↑	Δr	MAE
	Full SWA-MPPI	**0.0425**	–	**0.724**	–	**9.77**
1	No-MPPI (consensus only)	0.0493	+.007	0.650	−.074	11.02
2	Always-on MPPI	0.0422	−.000	0.729	+.005	9.74
3a	Utility: quadratic	0.0429	+.000	0.725	+.001	9.68
3b	Utility: linear	0.0444	+.002	0.706	−.018	10.12
4	Without utilitarian (N=3)	0.0610	+.019	0.441	−.283	11.09
5	No debiasing	0.0476	+.005	0.743	+.019	9.67

is unnecessary (Germany), and (iii) a deployment cost of one extra batched forward pass—the MPPI loop itself runs on cached scalars and is free at the wall-clock level.

What is doing the work? The combination of the consensus-shift and Binary-ARGS baselines gives an unusually crisp answer to the mechanism question. Uniform persona consensus (CS-MEAN, ARGS-UNIF) recovers B5 PersonaConsensus exactly, ruling out "it's just having multiple personas". WVS-salience-weighted re-scoring without asymmetric loss-sensitivity (ARGS-WVS) is strictly worse than uniform in 12/15 countries, ruling out "it's the WVS weighting". Magnitude clamping (CS-CLAMP) closes most of the gap, showing that *bounded* corrections are critical-and the residual 0.003 JSD that SWA-MPPI claws back over CS-CLAMP is fully recovered by replacing the clamp with the Prospect-Theory utility (ARGS-PT, JSD 0.0485, r 0.514, best/tied in 13/15 countries). The active ingredient is therefore the gain/loss asymmetry of Prospect Theory operating on top of bounded KL-regularised perturbations, which is exactly what the SWA-MPPI free-energy objective (Eq. 8) prescribes.

The Anglosphere pattern, treated as an open question. The largest gains occur for the English-speaking countries we evaluate—which are also the populations one would naively expect to be *best* served by an English-trained model. We resist the temptation to explain this post-hoc. We do not have an experimental handle that would let us attribute the pattern to RLHF, to the language of the prompt, to the AMCE characteristics of those specific populations, or to some interaction. We report it as a finding that deserves a dedicated follow-up study.

Cross-model failure patterns. Per-country results (Appendix M) reveal that failure modes are *model-specific*, not country-specific. On Llama-3.1-70B, the largest JSD gains shift to Vietnam (+65.8%) and Mexico (+66.4%), while Brazil (-23.3%) and Saudi Arabia (-10.7%) regress—a different pattern from Qwen2.5-72B, where Brazil improves (+6.7%) and Saudi Arabia is the only regression (-2.3%).

Vietnam deserves a specific caveat. On Llama-3.1-70B, Vietnam's vanilla Pearson r is -0.519, and after SWA-MPPI it is -0.135—closer to zero, but still negative. The headline JSD improvement is 65.8%, yet the 6-D rank ordering is still wrong. The most plausible reading is that the correction compresses Vietnam's AMCE vector toward the population mean, which happens to lie closer to the Vietnamese targets than Llama's vanilla output does. That is not direction-aware alignment in any strong sense. This is the cleanest in-paper instance of the JSD/r dissociation, and the same phenomenon appears at panel level on Mistral-Large. Mistral-Large shows a striking dissociation: JSD improves modestly (+5.8%) but Pearson r *decreases* for 11/15 countries (mean $\Delta r = -0.108$), indicating that the MPPI correction improves distributional distance while perturbing the rank ordering of AMCE dimensions. Qwen2.5-32B degrades in 7/15 countries, confirming the logit-concentration failure mode described in Section 5. These patterns suggest that SWA-MPPI's effectiveness depends on the base model's logit calibration quality, and that a pre-flight diagnostic (e.g., checking logit entropy) could predict when to activate the framework.

Limitations. We list the limitations we consider load-bearing, in roughly decreasing order of how much they should temper the reader's confidence.

Hyperparameter validation is the weakest link. Five of our hyperparameters (σ, λ_{coop}, α_{KL}, T_{dec}, T_{cat}) were tuned on 200 GPT-4-generated dilemmas with author-provided proxy labels. Our 50-scenario

MultiTP transfer check is genuinely too small to validate transfer of five hyperparameters; we report it for honesty, not as a guarantee. The temperature and dataset sweeps in §E–F are the strongest empirical defence we offer against this concern, and we encourage readers to weigh those more heavily than the synthetic-set tuning.

Single seed. As stated in §4.2, every result is from one run per model with one MPPI random draw per scenario; across-seed variance of δ^* is not characterised in this paper.

Pearson r in 6D is noisy. As noted in §5, individual r deltas should be read as directional rather than statistically separable. JSD is the metric our claims rest on.

Mistral-Large rank-order regression. On the largest model in our panel, JSD improves but Pearson r decreases for $11/15$ countries. We do not have a clean fix for this and it is consistent with the hypothesis that MPPI is partially compressing the AMCE vector toward the population mean on already-aligned models. A user deploying SWA-MPPI on a model whose vanilla alignment is already strong should evaluate per-dimension shifts before trusting the aggregate JSD.

Prospect Theory parameters are fixed a priori. We use the canonical $\alpha{=}\beta{=}0.88$, $\kappa{=}2.25$ from Kahneman and Tversky [1979] without per-country sensitivity analysis. These constants were estimated on Western experimental subjects and there is no a priori reason they should transfer cleanly to every culture in our panel; per-country or even learned PT parameters could in principle improve alignment, at the cost of new overfitting risks. We flag this as an open knob rather than a setting we have validated.

Utilitarian persona dominance. The ablation shows the utilitarian persona is the single most important component, which raises a fair question about whether it overwhelms the country-specific personas on non-numerosity dimensions. We do not currently disentangle utilitarian-driven shifts from WVS-driven shifts at the per-dimension level; the per-dimension breakdown discussed below is the natural place to test this.

No per-dimension AMCE breakdown. All headline metrics (JSD, Pearson r, MAE) aggregate over the six moral dimensions. We do not currently report which dimensions SWA-MPPI improves and which it hurts, and the per-category temperature T_{cat} (4.0 for Species, 3.5 for Gender, 1.5 for the others) raises a legitimate concern that the method may disproportionately correct Species at the expense of the other four. The utilitarian-persona ablation hints at a corresponding effect on the Utilitarianism dimension (Utilitarianism MPR drops from 61.2% to 48.0% when the persona is removed). A systematic per-dimension analysis is the single largest analytical gap in this version of the paper and is high on our follow-up list.

No human evaluation. We align to AMCE vectors derived from Moral Machine crowdsourcing. Whether AMCE-aligned outputs are perceived as more culturally appropriate by real humans is untested in this paper.

WVS-to-trolley linkage is indirect. The mapping from WVS value dimensions to trolley moral dimensions operates through what we call indirect modulation rather than a validated causal chain (Appendix B.3). The framework does not require an exact mapping—only that personas produce informative diversity in logit space—but we acknowledge this means we cannot make strong claims of the form "WVS dimension X caused AMCE shift Y."

Coverage and preprocessing. We evaluate 15 of 100+ MultiTP countries. A full zero-preprocessing evaluation is forthcoming (§5.4). WVS Wave 7 ($\sim$2017–2022) reflects a specific historical moment; persona quality is limited where Wave 7 coverage is sparse (e.g., Saudi Arabia, Brazil), which is visible in the per-country results. Extended limitations are in Appendix R.

Broader impact and safeguards. SWA-MPPI could improve fairness and inclusivity by calibrating LLM moral behaviour to diverse cultural preferences without costly retraining; the Anglosphere pattern we observe (§7) is itself a reminder that alignment protocols may need to be re-evaluated even for the populations they appear to serve best. However, the framework also risks aligning to harmful majority preferences (e.g., discriminatory norms against minorities within a population) or reifying stereotypes embedded in WVS survey averages. We identify three safeguards for responsible deployment: (i) *governance oversight*: SWA-MPPI should not override safety-critical model guardrails, and deployment decisions should involve domain experts familiar with the target culture; (ii) *periodic persona updates*: WVS data reflects a specific historical moment (Wave 7, $\sim$2017–2022),

and personas should be refreshed as newer waves become available to prevent temporal staleness; (iii) *robustness checks*: practitioners should audit per-dimension AMCE shifts to detect cases where persona-driven corrections entrench harmful norms (e.g., extreme gender biases in certain WVS profiles); and (iv) *explicit ethical bounds in the utility*: a natural extension of the SWA-MPPI utility is to clip δ^* on dimensions deemed sensitive (e.g., Gender, Social Value) either by a hard cap or by adding a steep penalty term to U_{total}, formalising what is currently a deployment-time audit into a decoding-time constraint. The WVS grounding provides a degree of empirical accountability-persona descriptions reflect surveyed population averages rather than author preferences-but this does not eliminate the risk of culturally legitimised harm.

8 Conclusion

We presented SWA-MPPI, a training-free inference-time framework for cross-cultural moral alignment. The story we want the reader to take away is narrower than the headline numbers might suggest. On strong instruction-tuned 70B-class models, persona disagreement combined with a Prospect-Theory-weighted MPPI correction in logit space reduces distance to country-level human moral preferences by a meaningful margin, and the method remains self-limiting when the personas already agree. On smaller quantised models the picture is less clean because vanilla JSD is already low; on Mistral-Large the picture is mixed because rank-order alignment can degrade even as JSD improves; on Qwen2.5-32B the method effectively does nothing, for a logit-level reason we diagnose explicitly. Our robustness suite is designed to rule out the confounders we could name—preprocessing, temperature, τ leakage, WVS weighting in isolation, deterministic shifts—and points consistently to Prospect-Theory asymmetry on KL-regularised perturbations as the load-bearing ingredient. Future directions include: adaptive σ selection based on per-model logit entropy for broader compatibility; extension from binary forced-choice to multi-option and open-ended moral reasoning, where a single decision logit is no longer available and the persona-disagreement signal must be lifted to a distribution over rationales; per-country or learned Prospect-Theory parameters; a per-dimension mediation analysis to validate the WVS-to-trolley mapping that this paper treats as a working assumption; integration with newer WVS waves to assess longitudinal drift; disentangling language from culture via multilingual evaluation within single countries; and formalising the ethical bounds described in the broader-impact paragraph as a decoding-time constraint inside U_{total}.

Acknowledgments and Disclosure of Funding

We thank the creators of MultiTP [Jin et al., 2025] and the Moral Machine experiment [Awad et al., 2018] for public data. Experiments on NVIDIA H100 hardware.

References

Aakanksha et al. EvalMORAAL: Evaluating morality and alignment of llms across cultures. *arXiv preprint arXiv:2410.01999*, 2024.

A. Arditi, O. Obeso, A. Syed, D. Paleka, N. Panickssery, W. Gurnee, and N. Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, volume 37, pages 136037–136083, 2025.

E. Awad, S. Dsouza, R. Kim, J. Schulz, J. Henrich, A. Shariff, J.-F. Bonnefon, and I. Rahwan. The moral machine experiment. *Nature*, 563(7729):59–64, 2018.

Y. Bai, S. Kadavath, S. Kundu, A. Askell, et al. Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*, 2022.

A. Deshpande, V. Murahari, T. Rajpurohit, A. Kalyan, and K. Narasimhan. Toxicity in ChatGPT: Analyzing persona-assigned language models. *arXiv preprint arXiv:2304.05335*, 2023.

E. Durmus, K. Nguyen, T. I. Liao, N. Schiefer, A. Askell, A. Bakhtin, et al. Towards measuring the representation of subjective global opinions in language models. *arXiv preprint arXiv:2306.16388*, 2023.

A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

C. Haerpfer, R. Inglehart, A. Moreno, C. Welzel, K. Kizilova, J. Diez-Medrano, M. Lagos, P. Norris, E. Ponarin, and B. Puranen. World Values Survey: Round seven – country-pooled datafile. Madrid, Spain & Vienna, Austria: JD Systems Institute & WVSA Secretariat, 2022. `https://doi.org/10.14281/18241.20`.

D. Hendrycks, C. Burns, S. Basart, A. Critch, J. Li, D. Song, and J. Steinhardt. Aligning AI with shared human values. In *International Conference on Learning Representations*, 2021.

J. Henrich, S. J. Heine, and A. Norenzayan. The weirdest people in the world? *Behavioral and Brain Sciences*, 33(2-3):61–83, 2010.

Z. Huang, W. Chen, C. Lyu, et al. MPPI-guided large language model generation. *arXiv preprint arXiv:2403.09269*, 2024.

Z. Jin, S. Levine, F. Gonzalez Adauto, O. Kamal, M. Sap, M. Sachan, R. Mihalcea, J. Tenenbaum, and B. Schölkopf. When to make exceptions: Exploring language models as accounts of human moral judgment. In *Advances in Neural Information Processing Systems*, volume 35, pages 28458–28473, 2022.

Z. Jin, M. Kleiman-Weiner, G. Piatti, S. Levine, J. Liu, F. Gonzalez, F. Ortu, A. Strausz, M. Sachan, R. Mihalcea, Y. Choi, and B. Schölkopf. Language model alignment in multilingual trolley problems. *arXiv preprint arXiv:2407.02273*, 2025. Published as a conference paper at ICLR 2025.

D. Kahneman and A. Tversky. Prospect theory: An analysis of decision under risk. *Econometrica*, 47(2):263–291, 1979.

M. Khanov, J. Burapacheep, and Y. Li. ARGS: Alignment as reward-guided search. *arXiv preprint arXiv:2402.01694*, 2024.

H. R. Kirk, A. Whitefield, P. Röttger, A. Bean, et al. The PRISM alignment project: What participatory, representative and individualised human feedback reveals about the subjective and multicultural alignment of large language models. *arXiv preprint arXiv:2404.16019*, 2024.

M. Kleiman-Weiner, R. Saxe, and J. B. Tenenbaum. Learning a commonsense moral theory. *Cognition*, 167:107–123, 2017.

Y. Li et al. Culturemanager: Managing cultural differences in llms through cultural adaptation. *arXiv preprint arXiv:2409.02945*, 2024.

N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, M. Bevilacqua, F. Petroni, and P. Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024.

Mistral AI Team. Mistral large 2 (instruct-2407). `https://mistral.ai/news/mistral-large-2407/`, 2024. Model card, accessed 2025.

S. Mudgal, J. Lee, H. Ganapathy, Y. Li, T. Wang, Y. Huang, Z. Chen, H.-T. Cheng, M. Collins, T. Strohman, et al. Controlled decoding from language models. *arXiv preprint arXiv:2310.17022*, 2024.

J. Myung et al. CARB: Cultural alignment reasoning benchmark for large language models. *arXiv preprint arXiv:2406.14766*, 2024.

A. Nie, Y. Zhang, A. Amdekar, C. Piech, T. B. Hashimoto, and T. Gerstenberg. MoCa: Measuring human-language model alignment on causal and moral judgment tasks. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

L. Ouyang, J. Wu, X. Jiang, D. Almeida, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.

Qwen Team. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2025.

M. J. Ryan, W. Held, and D. Yang. Unintended impacts of LLM alignment on global representation. *arXiv preprint arXiv:2402.15018*, 2024.

S. Santurkar, E. Durmus, F. Ladhak, C. Lee, P. Liang, and T. Hashimoto. Whose opinions do language models reflect? *arXiv preprint arXiv:2303.17548*, 2023.

T. Sorensen, J. Moore, J. Fisher, M. Gordon, N. Mireshghallah, C. M. Rytting, A. Ye, L. Jiang, X. Lu, N. Dziri, et al. A roadmap to pluralistic alignment. *arXiv preprint arXiv:2402.05070*, 2024.

Unsloth AI. Unsloth: Fast LLM fine-tuning and inference. `https://github.com/unslothai/unsloth`, 2024.

G. Williams, P. Drews, B. Goldfain, J. M. Rehg, and E. A. Theodorou. Information-theoretic model predictive control: Theory and applications to autonomous driving. *IEEE Transactions on Robotics*, 34(6):1603–1622, 2018.

J. Yao et al. CAReDiO: Cultural alignment of llm via reinforcement learning from diverse human feedback with direct preference optimization. *arXiv preprint arXiv:2410.18203*, 2024.

W. Zhang et al. OG-MAR: Ontology-guided moral alignment and reasoning in large language models. *arXiv preprint arXiv:2406.04264*, 2024.

A. Zou, L. Phan, S. Chen, J. Campbell, P. Guo, et al. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

A AMCE Estimation Details

Binary dimensions. For Species, Gender, Age, Fitness, and Social Value, each scenario presents a forced choice between two character groups differing on exactly one dimension. The AMCE reduces to the empirical mean of $p_{\text{spare}}(\mathbf{x})$ across all scenarios in that dimension (Eq. 1), consistent with MultiTP conventions [Jin et al., 2025].

Utilitarianism dimension. Utilitarianism scenarios vary the *number* of lives on each side rather than character identity. Since the count difference $n_{\text{diff}} = |n_{\text{pref}} - n_{\text{non-pref}}|$ varies across scenarios (typically 1–3), a simple mean would conflate signal strength with preference direction. We therefore fit an OLS regression $p_{\text{spare}} = a + b \cdot n_{\text{diff}}$ and evaluate at the mean observed n_{diff} in the data. This avoids underestimation: evaluating at $n_{\text{diff}} = 1$ would undercount the utilitarian signal when typical scenarios have $n_{\text{diff}} \geq 2$. Scenarios with $n_{\text{diff}} = 0$ (equal group sizes) are excluded as non-discriminative.

Human AMCE scale. The MultiTP ground-truth AMCEs in `country_specific_ACME.csv` are on a $[-1, 1]$ scale. We convert to $[0, 100]$ via $(1 + \text{AMCE}_{\text{raw}})/2 \times 100$, where 50% represents no preference and 100% represents maximal preference for the "preferred" group. Model AMCEs are computed on the same scale.

Consistency. Awad et al. [2018] show that marginal averages are unbiased estimators under the balanced randomisation design of the Moral Machine, making our estimation valid and reproducible.

B Persona Construction Details

B.1 Example Personas

The young-adult persona for the **United States**:

> *You are a young adult (18–35) from the United States of America. Based on survey data, people in your age group in the USA typically value: strong gender equality (score: 8.2/10), low religious centrality in moral decisions (score: 3.1/10), high personal autonomy (score: 7.8/10), high meritocratic orientation (score: 7.5/10), moderate family importance (score: 6.4/10), and moderate interpersonal trust (score: 5.3/10). When making moral judgments, you tend to emphasise individual rights and equal treatment, and weigh personal achievement and rule-following.*

The older-adult persona for **Vietnam**:

> *You are an older adult (55+) from Vietnam. Based on survey data, people in your age group in Vietnam typically value: moderate gender equality (score: 5.1/10), high religious and moral traditionalism (score: 7.2/10), strong family importance (score: 9.4/10), moderate personal autonomy (score: 4.8/10), and low interpersonal trust toward strangers (score: 3.2/10). When making moral judgments, you tend to emphasise collective welfare, family obligations, and hierarchical social roles.*

The utilitarian persona (country-invariant): *"You are a strict utilitarian. Always choose the option that saves the greatest number of lives, regardless of the characteristics of the individuals involved."*

B.2 WVS Data Processing Pipeline

We process WVS Wave 7 individual-level microdata as follows:

1. **Variable extraction:** For each respondent, we extract 8 value dimension scores from specific WVS items: gender egalitarianism (mean of Q58P, Q59P, Q60P), religious importance (Q6P), interpersonal trust (Q43P), moral permissiveness (mean of Q50, Q52P, Q54P), work centrality (Q5P), family importance (Q1P), personal autonomy (Q39P), meritocratic orientation (Q40P). All items use inverted scales where higher = more progressive/positive.
2. **Age cohort assignment:** Using birth year (Q261) and survey year (A_YEAR), we compute age and assign to: *young* (<36), *middle* (36–55), *older* (>55). Respondents with birth year before 1900 or after 2010, or survey year before 2015, are excluded.

3. **Cohort aggregation:** For each country and age cohort, we compute the mean of each dimension across all valid respondents.

4. **Natural-language conversion:** Dimension means are converted to descriptors using calibrated thresholds (e.g., religiosity ratio $>0.85 \rightarrow$ "deeply religious"; trust ratio $>0.55 \rightarrow$ "high interpersonal trust").

5. **Fallback:** For countries with insufficient WVS coverage (e.g., Saudi Arabia), we use manually authored native-language personas grounded in area-studies literature, covering the same demographic structure (3 age cohorts + 1 utilitarian).

B.3 WVS-to-Trolley Dimension Linkage

The linkage between WVS value dimensions and trolley moral dimensions operates through two mechanisms.

Direct thematic overlap: WVS *gender egalitarianism* maps onto the Gender dimension (societies scoring higher show weaker female-sparing preferences [Awad et al., 2018]); *religious importance* correlates with Species (human exceptionalism); *meritocratic orientation* relates to Social Value and Fitness.

Indirect modulation: Dimensions like *personal autonomy*, *interpersonal trust*, and *family importance* modulate the moral *frame*-e.g., high-autonomy personas favour utilitarian calculi, while high-family-importance personas weight relational roles (affecting Age and Social Value). Crucially, SWA-MPPI does not require exact causal mapping: the ensemble generates diversity in logit-space gaps, and MPPI selects the correction balancing collective utility. A rigorous per-dimension causal analysis (e.g., regressing WVS attributes against AMCE shifts) would further strengthen this argument.

C MPPI Theoretical Justification

C.1 Free-Energy Derivation of the KL Term

The scalar decision gap δ acts as a one-dimensional state and the perturbations ϵ_k as control inputs at a single horizon step. Under this view, the information-theoretic derivation of MPPI [Williams et al., 2018] applies directly: $w_k \propto \exp(U_{\text{total}}/\beta)$ minimises KL divergence between the sample distribution and the Boltzmann distribution of utility, so the KL term in Eq. 8 is a direct consequence of the free-energy objective rather than a heuristic penalty.

C.2 Comparison with Simpler Alternatives

Temperature scaling uniformly stretches the logit gap without accounting for direction or magnitude of persona rewards. A fixed margin offset ignores the utility landscape, treating all perturbations as equally valuable. MPPI performs an *adaptive, scenario-specific* correction via nonlinear importance weighting. The ablation (Table 8) provides indirect evidence: replacing PT with linear utility (closer to uniform weighting, hence a simple mean shift) degrades JSD by 4.5%.

D Hyperparameter Validation

Non-PT hyperparameters (σ, λ_{coop}, α_{KL}, T_{dec}, T_{cat}) were validated on a held-out synthetic set of 200 binary moral dilemmas generated by GPT-4, covering the same six dimensions as MultiTP with novel character descriptions. Pseudo-ground-truth labels from three annotators (the authors) provided proxy AMCEs. Grid search minimised mean JSD on Qwen2.5-72B. Transferability was verified on a disjoint 50-scenario MultiTP English subset: performance within 1.2% relative JSD of oracle-tuned configuration. T_{cat} values reflect empirical logit magnitude distributions: Species and Gender produce systematically larger gaps, requiring higher temperatures. Because this tuning pool is synthetic and author-annotated, residual bias is possible; we therefore report broad robustness sweeps (Appendix E, Appendix F) to show conclusions are stable beyond one tuned point.

E Temperature Sensitivity

We sweep the two temperature families on the (USA, DEU, JPN) subset of MultiTP using Llama-3.1-70B. Sweep A varies $T_{\text{dec}} \in \{0.10, 0.25, 0.50, 0.75, 1.0, 2.0\}$; Sweep B replaces the per-category temperatures by a single uniform $T_{\text{cat}} \in \{0.5, 1.0, 1.5, 2.0, 3.0, 4.0, 6.0\}$; Sweep C varies the "Others" bucket $T_{\text{cat}} \in \{0.75, 1.0, 1.5, 2.0, 3.0, 4.0\}$ while holding $T_{\text{cat}}[\text{Species}]=4.0$, $T_{\text{cat}}[\text{Gender}]=3.5$. All three sweeps yield JSD spans below 0.010 ($0.0074,\ 0.0073,\ 0.0077$ respectively), and in every sweep the SWA-MPPI default (†) is strictly *more conservative* than the empirically best value ($\Delta\text{JSD} \in \{-0.0061, -0.0021, -0.0037\}$). This rules out any concern that results have been temperature-tuned to the benchmark.

Table 9: Temperature sensitivity. Defaults (†) are strictly worse than the best setting in every sweep, providing a conservative lower bound.

Sweep	Setting	JSD ↓	Pearson r ↑	ΔJSD
A: T_{dec}	0.10	0.0502	0.662	+0.0013
	0.25	0.0491	0.658	+0.0002
	0.50†	0.0489	0.630	-
	0.75	0.0476	0.636	−0.0013
	1.00	0.0466	0.641	−0.0023
	2.00 (best)	0.0428	0.627	−0.0061
	JSD span		**0.0074**	
B: Uniform T_{cat}	0.5	0.0521	0.642	+0.0052
	1.0	0.0492	0.673	+0.0023
	1.5†	0.0469	0.700	-
	2.0	0.0461	0.707	−0.0008
	3.0	0.0454	0.707	−0.0015
	4.0	0.0451	0.701	−0.0018
	6.0 (best)	0.0448	0.684	−0.0021
	JSD span		**0.0073**	
C: Others T_{cat}	0.75	0.0526	0.566	+0.0039
	1.00	0.0512	0.589	+0.0025
	1.50†	0.0487	0.634	-
	2.00	0.0475	0.660	−0.0012
	3.00	0.0458	0.693	−0.0029
	4.00 (best)	0.0450	0.706	−0.0037
	JSD span		**0.0077**	

The monotonic JSD reduction with larger T_{dec} reflects the role of T_{dec} in undoing RLHF logit compression: higher values widen the effective token-probability landscape so that MPPI perturbations exert larger directional influence. However, r peaks near $T_{\text{dec}}=1.0$ and declines at 2.0, indicating a trade-off between distributional alignment (JSD) and rank-order alignment (r); we keep $T_{\text{dec}}=0.5$ as a deliberately conservative operating point.

F Dataset Preprocessing and Sensitivity

Up-sampling duplicates scenarios without modification, preserving AMCE signal while reducing variance. Category capping prevents dimension dominance. Side balancing enables effective debiasing. To evaluate whether these choices alter the effective evaluation distribution, we compare *six* preprocessing configurations on 10 representative countries (USA, DEU, CHN, JPN, BRA, VNM, GBR, KOR, RUS, NGA) using Llama-3.1-70B.

D1-NoAug shows the largest degradation ($+0.0025$), justifying synthetic augmentation of the under-represented Species and Utilitarianism categories. D4-NoFlip is marginally lower than D0, indicating that group-side randomisation is conservative (it adds noise) rather than harmful; we retain it because it preserves positional-debiasing integrity. The 0.0034 aggregate range is well below the ± 0.004 bootstrap CI, so reported results are not confounded by preprocessing choices.

Table 10: Per-country dataset sensitivity. JSD range = 0.0034 (aggregate) and < 0.010 in 9/10 countries; the only exception is RUS at 0.013, driven by augmentation sensitivity in the rare Species/Utilitarianism categories. D0 is retained as default.

Country	D0-Default	D1-NoAug	D2-Cap40	D3-Cap120	D4-NoFlip	D5-Strict
USA	0.0655	0.0674	0.0625	0.0644	0.0643	0.0684
DEU	0.0465	0.0515	0.0492	0.0467	0.0444	0.0484
CHN	0.0393	0.0363	0.0385	0.0383	0.0409	0.0398
JPN	0.0304	0.0291	0.0319	0.0296	0.0273	0.0320
BRA	0.0492	0.0551	0.0479	0.0510	0.0520	0.0480
VNM	0.0592	0.0574	0.0631	0.0585	0.0578	0.0626
GBR	0.0614	0.0667	0.0608	0.0618	0.0597	0.0633
KOR	0.0371	0.0354	0.0378	0.0357	0.0349	0.0382
RUS	0.0383	0.0489	0.0389	0.0374	0.0383	0.0357
NGA	0.0522	0.0568	0.0506	0.0525	0.0514	0.0538
Mean	**0.0479**	0.0505	0.0481	0.0476	0.0471	0.0490
Δ vs D0	-	$+0.0025$	$+0.0002$	-0.0003	-0.0008	$+0.0011$

G Held-Out τ Calibration

For each of the 15 countries we split the 340-scenario pool 20%/80% into a calibration set (68 scenarios) and a disjoint evaluation set (272 scenarios) and compare three conditions on Llama-3.1-70B: **(A)** τ calibrated on the held-out 20% slice and evaluated on the 80% slice (correct protocol; this is the configuration used in all main-paper experiments); **(B)** τ calibrated on the evaluation set itself (simulated leakage); and **(C)** fixed τ=0.001 (no calibration).

Table 11: Per-country held-out τ calibration. $|\Delta\text{JSD}|$ remains below the bootstrap CI (± 0.004) in all 15 countries; held-out τ is strictly more conservative (mean $\tau_B/\tau_A = 1.33$) yet alignment is unchanged.

Country	τ_A	τ_B	τ_B/τ_A	JSD_A	JSD_B	$\lvert\Delta\text{JSD}\rvert$	$A \leq B$
USA	0.332	0.586	1.77	0.0620	0.0637	0.0017	✓
DEU	0.809	0.496	0.61	0.0377	0.0387	0.0010	✓
CHN	0.250	0.285	1.14	0.0344	0.0346	0.0002	✓
JPN	0.824	1.109	1.35	0.0335	0.0334	0.0001	≈
BRA	1.563	1.961	1.25	0.0606	0.0647	0.0041	✓
SAU	0.773	1.305	1.69	0.0533	0.0534	0.0001	✓
VNM	0.131	0.247	1.89	0.0599	0.0600	0.0001	✓
GBR	0.289	0.411	1.42	0.0636	0.0646	0.0010	✓
KOR	0.210	0.196	0.93	0.0504	0.0504	0.0000	≈
RUS	0.091	0.110	1.21	0.0442	0.0442	0.0000	≈
MEX	0.355	0.520	1.46	0.0383	0.0387	0.0004	✓
NGA	0.922	0.891	0.97	0.0495	0.0491	0.0004	≈
AUS	0.285	0.389	1.37	0.0690	0.0704	0.0014	✓
FRA	0.906	1.555	1.72	0.0486	0.0485	0.0001	≈
IND	0.371	0.461	1.24	0.0479	0.0479	0.0000	≈
Mean	0.541	0.701	1.33	**0.0502**	0.0508	0.0007	12/15

The aggregate gap $|\text{JSD}_B - \text{JSD}_A| = 0.0006$ is smaller than the ± 0.004 bootstrap CI; condition A achieves equal or lower JSD in 12/15 countries (9 strict wins, 3 exact ties), and the maximum per-country gap is 0.0041 (BRA), still within noise. Notably $\tau_A < \tau_B$ in 12/15 countries (mean ratio $1.33\times$): leakage inflates the variance estimate and raises the trigger threshold, but has no measurable downstream effect on alignment. This robustness is theoretically expected because τ observes a single scalar (inter-agent logit variance) with no access to AMCE labels and targets a fixed 35% trigger rate rather than any alignment objective. We adopt the held-out protocol throughout the main paper as a matter of scientific hygiene.

H Cross-Lingual Token Elicitation

Native-language prompting system. Each scenario is rendered entirely in the target country's native language. The implementation includes full translations of: (i) scenario framing and context sentences (4 paraphrase variants per language), (ii) all 22 character types with singular/plural forms

in each language, (iii) lane labels and group identifiers in native script, and (iv) closing questions with language-appropriate conjunctions and grammar.

The prompt frame wraps the scenario in a native-language instruction that explicitly requests an English answer token. Each frame follows the pattern: "[moral dilemma preamble in native language] \n {scenario} \n [instruction to answer A or B in English]". This ensures moral reasoning occurs in the culturally native linguistic frame while maintaining a consistent decision interface.

Token verification. Token IDs are verified per model at initialisation by encoding the uppercase strings `"A"` and `"B"` with `add_special_tokens=False` and asserting single-token mappings: Qwen2.5: A=22397, B=11572; Llama-3.1: A=31266, B=9268. Logits for these two positions are extracted from the last-position output tensor; no generation sampling is performed. The model uses its built-in `chat_template` for correct role formatting across model families.

I Baseline Implementation Details

I.1 Activation Steering

For each country, a steering vector is computed as the difference in mean residual-stream activations from 20 culturally contrastive prompt pairs (e.g., "Answer as someone from [country] who values [WVS-high trait]" vs. "Answer as someone with the opposite values"). Vectors are extracted from layer 32 (transformer midpoint) of Llama-3.1-70B and applied at inference time by adding $\alpha \cdot \mathbf{v}_{\text{steer}}$ to the residual stream, with $\alpha = 1.5$ selected via the same synthetic validation set used for SWA-MPPI.

Sensitivity analysis and fairness considerations. Activation steering performance is known to be sensitive to (i) the extraction layer, (ii) the scaling coefficient α, and (iii) the construction of contrastive pairs [Arditi et al., 2025]. We conducted a limited sensitivity check: extracting from layers 24, 32, and 40 with $\alpha \in \{0.5, 1.0, 1.5, 2.0, 3.0\}$, we found that layer 32 with $\alpha = 1.5$ yielded the lowest mean JSD (0.0877); other configurations ranged from 0.089 to 0.112. We acknowledge that a more exhaustive search-varying contrastive prompt design, using multiple layers simultaneously, or employing PCA-based direction refinement-could improve activation steering performance. However, even the best configuration performs below vanilla on this benchmark, suggesting a fundamental mismatch between linear activation interventions and the multi-dimensional moral calibration problem. The key structural advantage of SWA-MPPI is that it operates *per-scenario* with adaptive correction magnitude, whereas activation steering applies a fixed direction uniformly across all scenarios.

I.2 ARGS and Controlled Decoding: from Incompatible to Reformulated

ARGS [Khanov et al., 2024] and controlled decoding [Mudgal et al., 2024] were originally formulated with trained reward models, which in the cross-cultural setting would demand a separate reward model per country-unavailable for MultiTP. However, the binary forced-choice structure of MultiTP admits a natural reformulation in which the reward is a closed-form functional of the persona logit gaps $\{\delta_i\}$. We exploit this in Appendix O (**Binary-ARGS**) to construct three reward specifications (ARGS-UNIF, ARGS-WVS, ARGS-PT) that share SWA-MPPI's batched forward pass and differ only in the reward kernel. The strongest variant, ARGS-PT, recovers the SWA-MPPI kernel exactly, providing an independent route to the same optimum.

J Detailed Ablation Analysis

Always-on vs. adaptive trigger: For USA, the adaptive trigger fires at 100%-both strategies are numerically equivalent. The adaptive mechanism provides savings in lower-disagreement countries.

Utilitarian persona: Removing it collapses Utilitarianism MPR from 61.2% to 48.0% (human: 76.6%), confirming it anchors the numerosity dimension.

Debiasing: Mean positional bias is 0.193 logit units; Species preference inflates from 85.8% to 93.5% without it (human: 79.2%).

K Trigger Mechanism and Latency

The adaptive τ targets 35% activation but achieves 100% for USA (all scenarios contested). The *flip rate* (MPPI reversals) is 6.0%, indicating confidence modulation rather than decision reversal.

Mean latency: 730ms (SWA-MPPI) vs. 160ms (vanilla) on Qwen2.5-72B (H100). Overhead from batched multi-agent passes and debiasing; MPPI operates on cached scalars and contributes negligibly.

L Germany Pearson r Analysis

Germany's JSD remains unchanged (0.0603) while Pearson r drops from 0.845 to 0.747. In 6D, Pearson r is sensitive to rank perturbations: a single dimension swap reduces r by $\sim$0.1 while JSD (depending on magnitudes, not ordering) stays constant. Rare MPPI corrections introduce minor rank perturbations without affecting distributional distance.

M Per-Country Results for All Models

Tables 12–16 present per-country results for all evaluated models beyond Qwen2.5-72B (Table 5 in the main text).

Table 12: Per-country results for **Llama-3.1-70B-Instruct**.

Country	Lang.	JSD ↓		Pearson r ↑		JSD Improv.	
		Van.	SWA	Van.	SWA	(Δ)	(%)
USA	en	.1561	.1234	0.012	0.197	−.033	+20.9
Germany	de	.0709	.0391	−0.455	0.685	−.032	+44.9
China	zh	.0459	.0311	0.406	0.802	−.015	+32.3
Japan	ja	.0457	.0416	0.496	0.466	−.004	+9.0
Brazil	pt	.0532	.0656	−0.370	0.348	+.012	−23.3
Saudi Arabia	ar	.0539	.0596	0.043	−0.210	+.006	−10.7
Vietnam	vi	.1546	.0529	−0.519	−0.135	−.102	+65.8
France	fr	.0500	.0516	0.250	0.394	+.002	−3.0
India	hi	.0442	.0337	−0.315	0.715	−.010	+23.6
South Korea	ko	.0893	.0436	−0.272	0.551	−.046	+51.2
Great Britain	en	.1568	.0918	0.012	0.286	−.065	+41.5
Russia	ru	.0449	.0413	0.539	0.637	−.004	+7.9
Mexico	es	.1265	.0425	−0.371	0.307	−.084	+66.4
Nigeria	en	.1582	.0988	0.109	0.497	−.059	+37.5
Australia	en	.1551	.1010	0.039	0.224	−.054	+34.9
Mean		**.0937**	**.0612**	−0.026	0.384	−.033	**+34.7**

N Consensus-Shift Baselines

To isolate the contribution of MPPI's importance-weighted sampling from the simpler benefit of using persona consensus at all, we implement four deterministic shift baselines applied directly to the scalar consensus gap $\bar{\delta}$, all sharing the same batched forward pass and two-pass debiasing as SWA-MPPI: CS-MEAN ($\delta_{\text{opt}} = \bar{\delta}$, equivalent to B5 PersonaConsensus), CS-STD ($\bar{\delta} \pm \sigma_{\text{agents}}$ in the consensus direction), CS-SCALED ($\bar{\delta}(1 + \alpha\,\text{Var})$), and CS-CLAMP ($\text{clip}(1.5\bar{\delta}, \pm 1)$).

CS-MEAN reproduces B5 PersonaConsensus exactly (cross-validating our ablation hierarchy). CS-STD *degrades* relative to CS-MEAN in 9/15 countries-variance amplification without KL regularisation actively hurts in high-disagreement regimes. CS-CLAMP is the strongest deterministic shift; its largest improvements occur in high inter-persona-variance countries (USA -0.015, AUS -0.016, GBR -0.012), exactly where Prospect Theory's asymmetric weighting is most active. The residual aggregate gap CS-CLAMP→SWA-MPPI (ΔJSD ≈ 0.003) is attributed to PT's gain/loss asymmetry, which no scalar shift can replicate.

Table 13: Per-country results for **Qwen2.5-7B-Instruct** (4-bit).

Country	Lang.	JSD ↓		Pearson r		JSD Improv.	
		Van.	SWA	Van.	SWA	(Δ)	(%)
USA	en	.0531	.0448	0.048	0.531	−.008	+15.5
Germany	de	.0530	.0522	0.207	0.212	−.001	+1.4
China	zh	.0512	.0455	0.401	0.397	−.006	+11.2
Japan	ja	.0527	.0540	0.490	0.449	+.001	−2.3
Brazil	pt	.0445	.0331	0.415	0.837	−.011	+25.8
Saudi Arabia	ar	.0539	.0443	0.234	0.265	−.010	+17.9
Vietnam	vi	.0547	.0472	0.259	0.263	−.008	+13.6
France	fr	.0529	.0442	0.259	0.315	−.009	+16.5
India	hi	.0570	.0524	0.466	0.183	−.005	+8.1
South Korea	ko	.0492	.0329	0.239	0.363	−.016	+33.2
Great Britain	en	.0520	.0302	0.140	0.483	−.022	+41.9
Russia	ru	.0506	.0452	0.373	0.491	−.005	+10.5
Mexico	es	.0479	.0384	0.038	0.360	−.010	+19.9
Nigeria	en	.0551	.0411	0.101	0.425	−.014	+25.3
Australia	en	.0550	.0431	0.170	0.452	−.012	+21.7
Mean		**.0522**	**.0426**	0.252	0.402	−.010	**+18.4**

Table 14: Per-country results for **Llama-3.1-8B-Instruct** (4-bit).

Country	Lang.	JSD ↓		Pearson r		JSD Improv.	
		Van.	SWA	Van.	SWA	(Δ)	(%)
USA	en	.0470	.0333	−0.323	0.498	−.014	+29.1
Germany	de	.0716	.0529	−0.101	0.382	−.019	+26.1
China	zh	.0561	.0457	−0.082	0.388	−.010	+18.5
Japan	ja	.0518	.0519	0.545	0.395	+.000	−0.2
Brazil	pt	.0732	.0628	−0.222	0.109	−.010	+14.2
Saudi Arabia	ar	.0551	.0513	−0.252	−0.230	−.004	+6.9
Vietnam	vi	.0505	.0423	−0.330	0.303	−.008	+16.2
France	fr	.0515	.0612	−0.610	−0.220	+.010	−18.8
India	hi	.0756	.0564	−0.471	0.347	−.019	+25.3
South Korea	ko	.0522	.0419	−0.352	−0.062	−.010	+19.7
Great Britain	en	.0375	.0255	−0.065	0.457	−.012	+32.0
Russia	ru	.0560	.0434	−0.230	0.593	−.013	+22.5
Mexico	es	.0439	.0383	−0.079	0.344	−.006	+12.7
Nigeria	en	.0474	.0410	−0.229	0.303	−.006	+13.5
Australia	en	.0505	.0363	−0.156	0.328	−.014	+28.1
Mean		**.0533**	**.0436**	−0.247	0.271	−.011	**+22.6**

O Binary-ARGS Reward-Guided Decoding

We adapt ARGS [Khanov et al., 2024] to the binary forced-choice setting as *Binary-ARGS*: K perturbations $\delta_k = \bar{\delta} + \epsilon_k$ ($\epsilon \sim \mathcal{N}(0, \sigma^2)$) are sampled and re-weighted by $\exp(r_k/\beta)$, where r_k is one of three closed-form rewards. ARGS-UNIF uses a uniform directional reward ($\text{sign}(\bar{\delta}) \cdot \delta_k/N$), equivalent to importance-weighted consensus averaging; ARGS-WVS weights each persona by its WVS cultural salience; ARGS-PT replaces the linear reward by the Prospect-Theory utility identical to the SWA-MPPI kernel.

Three findings emerge. **(1)** ARGS-UNIF (JSD 0.0524) closely matches B5 PersonaConsensus, confirming uniform importance weighting reduces to consensus averaging. **(2)** ARGS-WVS (JSD 0.0544) is *worse* than ARGS-UNIF in 12/15 countries: WVS salience weighting without asymmetric loss-sensitivity actively misaligns persona contributions. **(3)** ARGS-PT (JSD 0.0485, r 0.514) is best/tied in 13/15 countries and dominates on every aggregate metric, empirically showing that SWA-MPPI is precisely an ARGS variant with KL-regularised Prospect Theory. The two countries where ARGS-PT underperforms (BRA, IND) have the highest inter-persona variance and overcorrect without the KL term that SWA-MPPI provides through Eq. 8.

Table 15: Per-country results for **Mistral-Large-2407** (4-bit). Note: despite modest JSD gains, Pearson r decreases for 11/15 countries.

Country	Lang.	JSD ↓		Pearson r		JSD Improv.	
		Van.	SWA	Van.	SWA	(Δ)	(%)
USA	en	.0414	.0211	0.841	0.834	−.020	+49.0
Germany	de	.0378	.0291	0.734	0.675	−.009	+23.0
China	zh	.0572	.0618	0.824	0.585	+.005	−8.1
Japan	ja	.0457	.0535	0.814	0.387	+.008	−17.1
Brazil	pt	.0519	.0625	0.359	0.510	+.011	−20.4
Saudi Arabia	ar	.0238	.0315	0.724	0.671	+.008	−32.4
Vietnam	vi	.0521	.0428	0.768	0.571	−.009	+17.8
France	fr	.0481	.0475	0.683	0.628	−.001	+1.2
India	hi	.0457	.0527	0.772	0.631	+.007	−15.4
South Korea	ko	.0332	.0320	0.546	0.355	−.001	+3.6
Great Britain	en	.0418	.0331	0.853	0.730	−.009	+20.7
Russia	ru	.0552	.0566	0.768	0.631	+.001	−2.6
Mexico	es	.0472	.0481	0.748	0.642	+.001	−1.9
Nigeria	en	.0401	.0301	0.128	0.131	−.010	+25.0
Australia	en	.0516	.0338	0.865	0.682	−.018	+34.5
Mean		**.0457**	**.0430**	0.658	0.550	−.003	**+5.8**

Table 16: Per-country results for **Qwen2.5-32B-Instruct** (4-bit). SWA-MPPI degrades in 7/15 countries due to concentrated logit distributions.

Country	Lang.	JSD ↓		Pearson r		JSD Improv.	
		Van.	SWA	Van.	SWA	(Δ)	(%)
USA	en	.0447	.0440	0.887	0.738	−.001	+1.5
Germany	de	.0587	.0534	0.893	0.803	−.005	+9.1
China	zh	.0507	.0574	0.827	0.818	+.007	−13.1
Japan	ja	.0485	.0592	0.846	0.805	+.011	−22.0
Brazil	pt	.0717	.0645	0.552	0.422	−.007	+9.7
Saudi Arabia	ar	.0883	.1015	0.704	0.625	+.013	−14.9
Vietnam	vi	.0900	.0780	0.634	0.556	−.012	+13.1
France	fr	.0730	.0681	0.647	0.652	−.005	+6.7
India	hi	.0737	.0840	0.458	0.583	+.010	−13.8
South Korea	ko	.0529	.0303	0.106	0.064	−.023	+43.0
Great Britain	en	.0470	.0345	0.889	0.732	−.013	+26.5
Russia	ru	.0700	.0780	0.721	0.711	+.008	−11.3
Mexico	es	.0775	.0933	0.632	0.381	+.016	−20.3
Nigeria	en	.0641	.0512	0.124	0.156	−.013	+20.1
Australia	en	.0541	.0425	0.883	0.781	−.012	+21.4
Mean		**.0632**	**.0633**	0.625	0.550	+.000	**−0.2**

P Entropy-Aware Adaptive σ for Qwen2.5-32B

We study whether an entropy-aware adaptive σ can address the Qwen2.5-32B failure mode (Section 5). At each scenario we measure the Shannon entropy H of the top-50 token distribution at the decision position and set $\sigma = \sigma_{\text{base}} \cdot \text{clip}(H/H_{\text{ref}},\ 0.2,\ 5.0)$ with H_{ref}=2.5 nats and σ_{base}=0.30. We compare against a fixed-σ sweep on 8 countries (USA, DEU, CHN, JPN, BRA, VNM, GBR, KOR).

The mean entropy is $H = 0.521$ nats, far below H_{ref}, confirming concentrated logits. Adaptive σ self-scales to $\bar{\sigma} = 0.086$ ($r(H, \sigma)$=0.948), validating the diagnostic but collapsing onto the low-σ regime. The best fixed σ=0.80 reaches JSD 0.0599 vs. 0.0613 at default-a marginal absolute gain-while the per-country breakdown exposes the deeper problem: countries cluster into a high-σ group (USA/JPN/GBR/KOR) and a low-σ group (DEU/CHN/BRA/VNM), but *all* countries share $H \approx 0.52$ nats, so a per-scenario entropy signal cannot discriminate them. A complete fix therefore requires *jointly* adaptive τ per country to detect which countries have sufficient logit variance to benefit from large perturbations-left as future work.

Table 17: Per-country JSD for the four consensus-shift baselines on Llama-3.1-70B (15 countries). CS-CLAMP is the strongest deterministic shift but still loses to SWA-MPPI in aggregate; CS-STD is catastrophic for NGA ($+0.017$ vs. CS-MEAN) due to unbounded variance amplification.

Country	CS-Mean	CS-Std	CS-Scaled	CS-Clamp
USA	0.0923	0.0919	0.0933	**0.0771**
DEU	**0.0410**	0.0448	0.0449	0.0428
CHN	0.0591	0.0618	0.0608	**0.0520**
JPN	**0.0306**	0.0335	0.0343	0.0311
BRA	0.0524	0.0519	0.0530	**0.0459**
SAU	0.0565	0.0584	0.0570	**0.0536**
VNM	0.0632	0.0659	0.0669	**0.0564**
FRA	0.0547	0.0516	0.0552	**0.0442**
IND	0.0421	**0.0410**	0.0413	0.0418
KOR	0.0451	0.0454	0.0451	**0.0443**
GBR	0.0800	0.0842	0.0844	**0.0681**
RUS	0.0450	**0.0435**	0.0443	0.0429
MEX	0.0539	0.0539	0.0531	**0.0472**
NGA	0.0826	0.0998	0.0987	**0.0799**
AUS	0.0868	0.0864	0.0861	**0.0711**
Mean	0.0590	0.0609	0.0612	**0.0532**
SWA-MPPI	**0.0502** (Pearson r = 0.424)			

Table 18: Binary-ARGS per-country results (Llama-3.1-70B, 15 countries). ARGS-PT achieves best or tied JSD in 13/15 countries and best r in 11/15, confirming SWA-MPPI is precisely an ARGS variant with KL-regularised Prospect-Theory utility.

	JSD ↓			Pearson r ↑		
Country	Unif	WVS	PT	Unif	WVS	PT
USA	0.0757	0.0793	**0.0646**	0.506	0.484	**0.626**
DEU	0.0476	0.0512	**0.0465**	0.668	0.656	**0.732**
CHN	0.0347	0.0396	0.0362	0.753	0.721	**0.788**
JPN	0.0308	0.0311	**0.0275**	0.782	0.757	**0.831**
BRA	**0.0482**	0.0509	0.0519	0.351	0.317	**0.375**
SAU	0.0541	0.0539	0.0534	0.076	0.079	**0.165**
VNM	0.0717	0.0755	**0.0551**	-0.507	-0.427	-0.369
GBR	0.0692	0.0709	**0.0588**	0.494	0.488	**0.616**
KOR	0.0421	0.0386	0.0364	0.722	**0.750**	0.748
RUS	**0.0399**	0.0443	0.0415	0.723	0.728	**0.782**
MEX	0.0464	0.0496	**0.0428**	0.255	0.256	**0.503**
NGA	0.0567	0.0598	**0.0502**	0.775	0.779	**0.852**
AUS	0.0742	0.0763	**0.0664**	0.461	0.452	**0.557**
FRA	0.0495	0.0484	0.0480	0.336	0.370	**0.423**
IND	**0.0457**	0.0463	0.0475	0.045	-0.003	**0.079**
Mean	0.0524	0.0544	**0.0485**	0.429	0.427	**0.514**

Q Extended Country Pipeline Run (Engineering Check)

We run the SWA-MPPI pipeline end-to-end on 8 additional countries not in the MultiTP benchmark, spanning 4 script systems and 5 language families: Poland (pl), Turkey (tr), Indonesia (id), Argentina (es), Ukraine (uk), Taiwan (zh), the Philippines (tl), and Egypt (ar). **This appendix makes an engineering claim, not a scientific one.** Because no Moral Machine human AMCEs are available for these countries, we evaluate on 300 synthetic scenarios per country whose Vanilla AMCE is uniformly 50.0 by construction. The only thing this experiment establishes is that the persona-construction, translation, batched-forward-pass, τ-calibration, and MPPI-update code runs to completion across the 4 scripts and 5 language families above. The $\sim+0.05$pp "AMCE shift" is therefore measuring perturbation activity, not alignment, and the mean JSD numbers should not be compared to the main table. We retain the appendix only as evidence that the framework is not silently English-specific.

Because the synthetic scenarios uniformly initialise Vanilla AMCE at 50.0, the shift metric measures only the pipeline's ability to modulate output distributions away from an uninformative prior; it cannot be interpreted as alignment improvement over human preferences and is not directly comparable to the main 15-country results. We include this probe purely to demonstrate that the WVS-grounded persona construction, the batched forward pass, the two-pass debiasing, and the MPPI loop all run end-to-end across right-to-left Arabic, CJK, Cyrillic, and several Latin-script languages without any

Table 19: Per-country JSD under fixed-σ on Qwen2.5-32B. Two latent country clusters emerge: USA/JPN/GBR/KOR strongly prefer high σ (USA: $0.0523 \rightarrow 0.0317$ at σ=1.0); DEU/CHN/BRA/VNM strictly prefer low σ. All 8 countries share $H \approx 0.52$ nats, so the entropy signal alone cannot discriminate them.

σ	USA	DEU	CHN	JPN	BRA	VNM	GBR	KOR	Mean
0.05	0.0515	0.0647	0.0533	0.0690	0.0692	0.0710	0.0500	0.0670	0.0620
0.10	0.0519	0.0648	0.0532	0.0684	0.0694	0.0711	0.0499	0.0668	0.0619
0.15	0.0522	0.0649	0.0530	0.0676	0.0698	0.0711	0.0499	0.0665	0.0619
0.20	0.0518	0.0651	0.0527	0.0661	0.0702	0.0713	0.0498	0.0661	0.0616
$0.30^{\dagger}$	0.0523	0.0654	0.0516	0.0644	0.0710	0.0717	0.0491	0.0649	0.0613
0.50	0.0487	0.0661	0.0515	0.0601	0.0732	0.0734	0.0476	0.0604	0.0601
0.80	**0.0427**	0.0672	0.0599	**0.0585**	0.0780	0.0757	**0.0439**	**0.0532**	**0.0599**
1.00	0.0317	0.0681	0.0714	0.0637	0.0833	0.0744	0.0458	0.0522	0.0613
Adaptive ($\bar{\sigma}$=0.086)	0.0516	0.0647	0.0532	0.0688	0.0693	0.0711	0.0500	0.0670	0.0620

Table 20: Extended country pipeline run (Llama-3.1-70B). The pipeline executes end-to-end for 4 script systems and 5 language families; MPPI-induced AMCE shifts are detected in all 8 countries (mean $+0.05 \pm 0.02$ pp).

Country	Lang.	Script	AMCE Shift	Vanilla JSD	SWA JSD
Poland (POL)	Polish	Latin	+0.03 pp	0.0496	0.0497
Turkey (TUR)	Turkish	Latin	+0.04 pp	0.0376	0.0377
Indonesia (IDN)	Indonesian	Latin	+0.06 pp	0.0425	0.0424
Argentina (ARG)	Spanish	Latin	+0.04 pp	0.0388	0.0387
Ukraine (UKR)	Ukrainian	Cyrillic	+0.07 pp	0.0485	0.0483
Taiwan (TWN)	Chinese	CJK	+0.03 pp	0.0427	0.0426
Philippines (PHL)	Filipino	Latin	+0.06 pp	0.0396	0.0396
Egypt (EGY)	Arabic	Arabic	+0.06 pp	0.0511	0.0513
Mean			$\mathbf{+0.05 \pm 0.02}$ **pp**	0.0438	0.0438

code change. Future work should collect human AMCE annotations for these countries to enable full quantitative evaluation.

R Extended Limitations

Inter-persona reward assumption. The reward $r_i = \delta_i - \delta_{\text{base}}$ assumes persona shifts approximate human target directions. This is plausible when the base model represents a WEIRD-biased prior, but if a persona's shift is orthogonal to the human target, the signal may mislead. Per-dimension validation is needed.

Country coverage. 15 countries across major cultural clusters, but no low-resource languages (Amharic, Khmer, Yoruba). Cross-lingual generalisation to such settings remains untested.

Binary format. The logit-difference formulation requires generalisation to softmax for multi-option settings.

Language–culture conflation. One language per country may conflate linguistic and cultural effects.

Model quantisation. 4-bit/Int8 quantisation may introduce logit artefacts affecting MPPI stability.

WVS temporal staleness. Wave 7 ($\sim$2017–2022) data may not reflect current preferences in rapidly evolving societies.

Qwen2.5-32B failure mode. Concentrated logit distributions (H=0.521 nats) cause MPPI perturbations to land off-manifold and yield near-uniform importance weights. Appendix P shows that an entropy-aware σ correctly diagnoses the cause but cannot fix it on its own, because all 8 evaluated

countries share the same low entropy and thus cannot be discriminated by a per-scenario entropy signal. A complete remedy requires *jointly* adaptive (σ, τ) per country and is left to future work.